

baseline

April 24, 2025

```
[1]: import os
import itertools
import random
import data_pipeline as pipeline
import pandas as pd
import numpy as np
import tensorflow as tf
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
from tensorflow.keras.metrics import BinaryAccuracy, Precision, Recall
from tensorflow.keras.callbacks import EarlyStopping
```

```
[2]: PROCESSED_DIR = "../data/processed"
IMG_HEIGHT = IMG_WIDTH = 224
NUM_CHANNELS = 3
NUM_CLASSES = 2
EPOCHS = 30
BATCH_SIZE = 128
```

```
[3]: SEED = 42
np.random.seed(SEED)
tf.random.set_seed(SEED)
random.seed(SEED)
```

```
[4]: metrics = [
    BinaryAccuracy(name="accuracy"),
    Precision(name="precision"),
    Recall(name="recall"),
]
```

```
[5]: train_dir = os.path.join(PROCESSED_DIR, "train")
val_dir = os.path.join(PROCESSED_DIR, "val")
test_dir = os.path.join(PROCESSED_DIR, "test")

all_paths = pipeline.get_image_paths(PROCESSED_DIR)
train_paths = [path for path in all_paths if "/train/" in path]
mean, std = pipeline.calc_mean_std(train_paths)
```

```

# use generators from data_pipeline for training, validation, and testing
print("loading train/val/test generators from data_pipeline")
train_data_gen, val_data_gen, test_data_gen, test_data_gen_raw = pipeline.
    ↪load_data(
        train_dir, val_dir, test_dir, mean, std
    )

```

```

loading train/val/test generators from data_pipeline
creating train generator
Found 1600 images belonging to 2 classes.
creating validation generator
Found 400 images belonging to 2 classes.
creating test generator (normalized)
Found 200 images belonging to 2 classes.
creating test generator (raw)
Found 200 images belonging to 2 classes.

```

```

[6]: image_size = (IMG_HEIGHT, IMG_WIDTH)
     input_shape = (IMG_HEIGHT, IMG_WIDTH, NUM_CHANNELS)

```

```

[7]: class_names = list(train_data_gen.class_indices.keys())
     print(f"class names found: {class_names}")

```

```

class names found: ['NORMAL', 'COVID']

```

```

[8]: def build_model(input_shape):
     """
     build keras sequential model

     params
     -----
     input_shape: tuple
         shape of input images (height, width, channels)

     returns
     -----
     model: tf.keras.Model
         compiled keras model
     """
     model = tf.keras.Sequential(
         [
             tf.keras.layers.Input(shape=input_shape),
             # convolutional
             tf.keras.layers.Conv2D(32, (3, 3), activation="relu", ↪
             ↪padding="same"),
             tf.keras.layers.MaxPooling2D((2, 2)),

```

```

        tf.keras.layers.Conv2D(64, (3, 3), activation="relu",
        ↪padding="same"),
        tf.keras.layers.MaxPooling2D((2, 2)),
        tf.keras.layers.Conv2D(128, (3, 3), activation="relu",
        ↪padding="same"),
        tf.keras.layers.MaxPooling2D((2, 2)),
        # fully connected
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(128, activation="relu"),
        tf.keras.layers.Dropout(0.3), # random value
        tf.keras.layers.Dense(64, activation="relu"),
        tf.keras.layers.Dropout(0.3), # random value
        tf.keras.layers.Dense(1, activation="sigmoid"),
    ]
)
return model

```

```

[9]: def train_model(model, train_data, val_data, epochs=EPOCHS):
    # compile the model
    model.compile(
        optimizer="adam",
        loss="binary_crossentropy",
        metrics=metrics,
    )
    # train the model
    history = model.fit(train_data, epochs=epochs, validation_data=val_data,
    ↪verbose=2)
    return history

```

```

[10]: # build model
model = build_model(input_shape)

# print model summary
print("model architecture:")
model.summary()

```

model architecture:

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 224, 224, 32)	896
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	18,496

max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
conv2d_2 (Conv2D)	(None, 56, 56, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 128)	0
flatten (Flatten)	(None, 100352)	0
dense (Dense)	(None, 128)	12,845,184
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8,256
dropout_1 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 1)	65

Total params: 12,946,753 (49.39 MB)

Trainable params: 12,946,753 (49.39 MB)

Non-trainable params: 0 (0.00 B)

```
[11]: # train the model
history = train_model(model, train_data_gen, val_data_gen, EPOCHS)
```

```
/opt/anaconda3/envs/ml-2025/lib/python3.12/site-
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
  self._warn_if_super_not_called()

Epoch 1/30
13/13 - 27s - 2s/step - accuracy: 0.6137 - loss: 0.7341 - precision: 0.6338 -
recall: 0.5387 - val_accuracy: 0.6000 - val_loss: 0.7182 - val_precision: 0.5625
- val_recall: 0.9000
Epoch 2/30
13/13 - 28s - 2s/step - accuracy: 0.7206 - loss: 0.5636 - precision: 0.7182 -
recall: 0.7262 - val_accuracy: 0.6225 - val_loss: 0.6532 - val_precision: 0.6244
- val_recall: 0.6150
Epoch 3/30
```

13/13 - 25s - 2s/step - accuracy: 0.7494 - loss: 0.5258 - precision: 0.7264 -
recall: 0.8000 - val_accuracy: 0.6075 - val_loss: 0.6518 - val_precision: 0.6972
- val_recall: 0.3800
Epoch 4/30
13/13 - 26s - 2s/step - accuracy: 0.7706 - loss: 0.4954 - precision: 0.7580 -
recall: 0.7950 - val_accuracy: 0.6900 - val_loss: 0.5980 - val_precision: 0.6583
- val_recall: 0.7900
Epoch 5/30
13/13 - 26s - 2s/step - accuracy: 0.7812 - loss: 0.4706 - precision: 0.7635 -
recall: 0.8150 - val_accuracy: 0.7775 - val_loss: 0.5133 - val_precision: 0.7211
- val_recall: 0.9050
Epoch 6/30
13/13 - 29s - 2s/step - accuracy: 0.7962 - loss: 0.4476 - precision: 0.7849 -
recall: 0.8163 - val_accuracy: 0.7550 - val_loss: 0.4970 - val_precision: 0.7217
- val_recall: 0.8300
Epoch 7/30
13/13 - 27s - 2s/step - accuracy: 0.8050 - loss: 0.4185 - precision: 0.7779 -
recall: 0.8537 - val_accuracy: 0.7775 - val_loss: 0.4781 - val_precision: 0.8284
- val_recall: 0.7000
Epoch 8/30
13/13 - 25s - 2s/step - accuracy: 0.8263 - loss: 0.3832 - precision: 0.7919 -
recall: 0.8850 - val_accuracy: 0.7875 - val_loss: 0.4408 - val_precision: 0.8010
- val_recall: 0.7650
Epoch 9/30
13/13 - 23s - 2s/step - accuracy: 0.8425 - loss: 0.3590 - precision: 0.8301 -
recall: 0.8612 - val_accuracy: 0.8275 - val_loss: 0.4040 - val_precision: 0.7885
- val_recall: 0.8950
Epoch 10/30
13/13 - 23s - 2s/step - accuracy: 0.8288 - loss: 0.3618 - precision: 0.8153 -
recall: 0.8500 - val_accuracy: 0.8275 - val_loss: 0.3938 - val_precision: 0.7835
- val_recall: 0.9050
Epoch 11/30
13/13 - 23s - 2s/step - accuracy: 0.8431 - loss: 0.3427 - precision: 0.8109 -
recall: 0.8950 - val_accuracy: 0.8075 - val_loss: 0.4080 - val_precision: 0.8639
- val_recall: 0.7300
Epoch 12/30
13/13 - 23s - 2s/step - accuracy: 0.8562 - loss: 0.3260 - precision: 0.8401 -
recall: 0.8800 - val_accuracy: 0.8275 - val_loss: 0.3707 - val_precision: 0.7860
- val_recall: 0.9000
Epoch 13/30
13/13 - 23s - 2s/step - accuracy: 0.8712 - loss: 0.3042 - precision: 0.8536 -
recall: 0.8963 - val_accuracy: 0.8775 - val_loss: 0.3510 - val_precision: 0.8719
- val_recall: 0.8850
Epoch 14/30
13/13 - 24s - 2s/step - accuracy: 0.8744 - loss: 0.2972 - precision: 0.8630 -
recall: 0.8900 - val_accuracy: 0.8700 - val_loss: 0.3299 - val_precision: 0.8524
- val_recall: 0.8950
Epoch 15/30

13/13 - 24s - 2s/step - accuracy: 0.8775 - loss: 0.2864 - precision: 0.8639 -
recall: 0.8963 - val_accuracy: 0.8600 - val_loss: 0.3379 - val_precision: 0.8364
- val_recall: 0.8950

Epoch 16/30

13/13 - 23s - 2s/step - accuracy: 0.8712 - loss: 0.3087 - precision: 0.8685 -
recall: 0.8750 - val_accuracy: 0.8700 - val_loss: 0.3494 - val_precision: 0.8627
- val_recall: 0.8800

Epoch 17/30

13/13 - 23s - 2s/step - accuracy: 0.8744 - loss: 0.3022 - precision: 0.8613 -
recall: 0.8925 - val_accuracy: 0.8800 - val_loss: 0.3266 - val_precision: 0.8762
- val_recall: 0.8850

Epoch 18/30

13/13 - 23s - 2s/step - accuracy: 0.8731 - loss: 0.3069 - precision: 0.8583 -
recall: 0.8938 - val_accuracy: 0.8725 - val_loss: 0.3273 - val_precision: 0.8821
- val_recall: 0.8600

Epoch 19/30

13/13 - 23s - 2s/step - accuracy: 0.8694 - loss: 0.3047 - precision: 0.8680 -
recall: 0.8712 - val_accuracy: 0.8650 - val_loss: 0.3511 - val_precision: 0.8763
- val_recall: 0.8500

Epoch 20/30

13/13 - 23s - 2s/step - accuracy: 0.8769 - loss: 0.2972 - precision: 0.8577 -
recall: 0.9038 - val_accuracy: 0.8800 - val_loss: 0.3117 - val_precision: 0.8838
- val_recall: 0.8750

Epoch 21/30

13/13 - 23s - 2s/step - accuracy: 0.8719 - loss: 0.2888 - precision: 0.8529 -
recall: 0.8988 - val_accuracy: 0.8800 - val_loss: 0.3061 - val_precision: 0.9043
- val_recall: 0.8500

Epoch 22/30

13/13 - 23s - 2s/step - accuracy: 0.8744 - loss: 0.2734 - precision: 0.8622 -
recall: 0.8913 - val_accuracy: 0.8750 - val_loss: 0.3278 - val_precision: 0.8827
- val_recall: 0.8650

Epoch 23/30

13/13 - 23s - 2s/step - accuracy: 0.8925 - loss: 0.2642 - precision: 0.8905 -
recall: 0.8950 - val_accuracy: 0.9000 - val_loss: 0.2958 - val_precision: 0.8960
- val_recall: 0.9050

Epoch 24/30

13/13 - 23s - 2s/step - accuracy: 0.8906 - loss: 0.2730 - precision: 0.8788 -
recall: 0.9062 - val_accuracy: 0.8900 - val_loss: 0.2887 - val_precision: 0.8611
- val_recall: 0.9300

Epoch 25/30

13/13 - 23s - 2s/step - accuracy: 0.8888 - loss: 0.2561 - precision: 0.8821 -
recall: 0.8975 - val_accuracy: 0.8750 - val_loss: 0.3221 - val_precision: 0.8906
- val_recall: 0.8550

Epoch 26/30

13/13 - 23s - 2s/step - accuracy: 0.8938 - loss: 0.2600 - precision: 0.8899 -
recall: 0.8988 - val_accuracy: 0.8875 - val_loss: 0.3082 - val_precision: 0.8974
- val_recall: 0.8750

Epoch 27/30

13/13 - 23s - 2s/step - accuracy: 0.9006 - loss: 0.2366 - precision: 0.8820 -
recall: 0.9250 - val_accuracy: 0.8775 - val_loss: 0.3197 - val_precision: 0.8995
- val_recall: 0.8500

Epoch 28/30

13/13 - 23s - 2s/step - accuracy: 0.9044 - loss: 0.2411 - precision: 0.8999 -
recall: 0.9100 - val_accuracy: 0.8700 - val_loss: 0.3466 - val_precision: 0.8776
- val_recall: 0.8600

Epoch 29/30

13/13 - 23s - 2s/step - accuracy: 0.9000 - loss: 0.2360 - precision: 0.8912 -
recall: 0.9112 - val_accuracy: 0.8850 - val_loss: 0.2876 - val_precision: 0.9053
- val_recall: 0.8600

Epoch 30/30

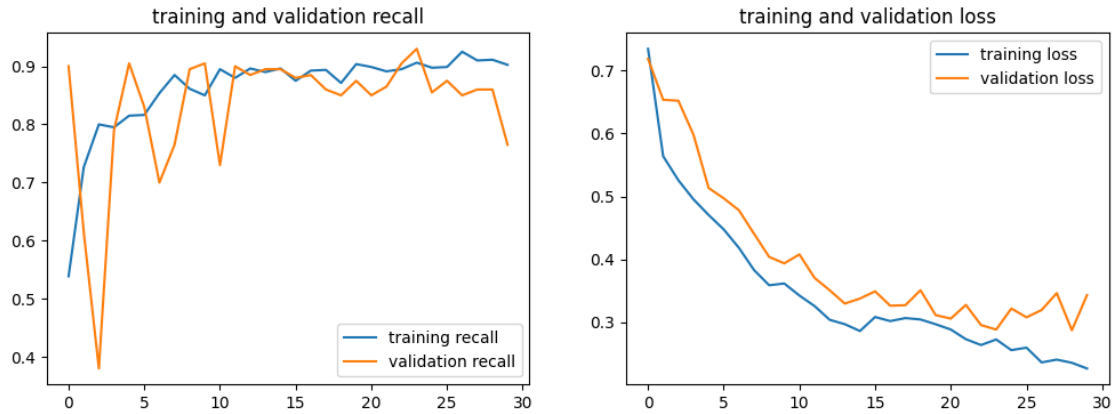
13/13 - 23s - 2s/step - accuracy: 0.8975 - loss: 0.2269 - precision: 0.8936 -
recall: 0.9025 - val_accuracy: 0.8550 - val_loss: 0.3434 - val_precision: 0.9329
- val_recall: 0.7650

```
[12]: # plot training history
recall = history.history["recall"]
val_recall = history.history["val_recall"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]

epochs_range = range(EPOCHS)

plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, recall, label="training recall")
plt.plot(epochs_range, val_recall, label="validation recall")
plt.legend(loc="lower right")
plt.title("training and validation recall")

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label="training loss")
plt.plot(epochs_range, val_loss, label="validation loss")
plt.legend(loc="upper right")
plt.title("training and validation loss")
plt.show()
```



```
[13]: # evaluate the model on the validation set after training
print("\nevaluating model on validation data after training")
results = model.evaluate(val_data_gen, verbose=1)
print(f"final validation loss: {results[0]}")
print(f"final validation accuracy: {results[1]}")
print(f"final validation precision: {results[2]}")
print(f"final validation recall: {results[3]}")
```

```
evaluating model on validation data after training
4/4          2s 355ms/step -
accuracy: 0.8876 - loss: 0.2720 - precision: 0.7192 - recall: 0.6218
final validation loss: 0.3434081971645355
final validation accuracy: 0.8550000190734863
final validation precision: 0.9329268336296082
final validation recall: 0.7649999856948853
```

1 with hyperparam tuning

```
[14]: def build_model(input_shape, filters_conv1, units_dense1, dropout_rate):
    """
    build keras sequential model

    params
    -----
    input_shape: tuple
        shape of input images (height, width, channels)
    filters_conv1: int
        number of filters in the first convolutional layer
    units_dense1: int
        number of units in the first dense layer
    dropout_rate: float
```



```

        dropout rate for dropout layers

    returns
    -----
    model: tf.keras.Model
        keras model (not compiled)
    """
    model = tf.keras.Sequential(
        [
            tf.keras.layers.Input(shape=input_shape),
            # convolutional
            tf.keras.layers.Conv2D(
                filters_conv1, (3, 3), activation="relu", padding="same"
            ),
            tf.keras.layers.MaxPooling2D((2, 2)),
            tf.keras.layers.Conv2D(64, (3, 3), activation="relu",
↳padding="same"),
            tf.keras.layers.MaxPooling2D((2, 2)),
            tf.keras.layers.Conv2D(128, (3, 3), activation="relu",
↳padding="same"),
            tf.keras.layers.MaxPooling2D((2, 2)),
            # fully connected
            tf.keras.layers.Flatten(),
            tf.keras.layers.Dense(units_dense1, activation="relu"),
            tf.keras.layers.Dropout(dropout_rate),
            tf.keras.layers.Dense(64, activation="relu"),
            tf.keras.layers.Dropout(dropout_rate),
            tf.keras.layers.Dense(1, activation="sigmoid"),
        ]
    )
    return model

```

```

[15]: # define search space
learning_rates = [1e-3, 1e-4, 1e-5]
filters_conv1_list = [16, 32, 64]
units_dense1_list = [64, 128, 256]
dropout_rates = [0.2, 0.3, 0.4, 0.5]

```

```

[16]: # create all possible combinations (full grid)
param_grid = list(
    itertools.product(
        learning_rates, filters_conv1_list, units_dense1_list, dropout_rates
    )
)

```

```

[17]: # define number of combinations to randomly sample
num_combinations_to_test = 5

```

```

# randomly sample combinations
sampled_params = random.sample(param_grid, num_combinations_to_test)
print(
    f"randomly sampling {num_combinations_to_test} combinations from
    ↳{len(param_grid)} total."
)

```

randomly sampling 5 combinations from 108 total.

```

[18]: # store results
results_list = []

```

```

[19]: # early stopping callback
early_stopping = EarlyStopping(
    monitor="val_loss", patience=5, restore_best_weights=True, verbose=1
)

```

```

[20]: print("starting hyperparameter tuning")

for lr, filters1, units1, dr in sampled_params:
    print(
        f"testing: lr={lr}, filters_conv1={filters1}, units_dense1={units1},
        ↳dropout_rate={dr}"
    )

    # build model
    input_shape = (IMG_HEIGHT, IMG_WIDTH, NUM_CHANNELS)
    model = build_model(
        input_shape,
        filters_conv1=filters1,
        units_dense1=units1,
        dropout_rate=dr,
    )

    # compile model
    optimizer = tf.keras.optimizers.Adam(learning_rate=lr)
    model.compile(optimizer=optimizer, loss="binary_crossentropy",
    ↳metrics=metrics)

    # train the model
    history = model.fit(
        train_data_gen,
        epochs=EPOCHS,
        validation_data=val_data_gen,
        callbacks=[early_stopping],
        verbose=2,
    )

```

```

)

# evaluate the model on the validation set using best weights from early_
↪stopping
print("evaluating best model from this run")
eval_results = model.evaluate(val_data_gen, verbose=1)

run_results = {
    "learning_rate": lr,
    "filters_conv1": filters1,
    "units_dense1": units1,
    "dropout_rate": dr,
    "val_loss": eval_results[0],
    "val_accuracy": eval_results[1],
    "val_precision": eval_results[2],
    "val_recall": eval_results[3],
    "epochs_trained": len(history.epoch),
}
results_list.append(run_results)

print("hyperparameter tuning finished.")

```

starting hyperparameter tuning

testing: lr=0.001, filters_conv1=16, units_dense1=64, dropout_rate=0.2

Epoch 1/30

13/13 - 19s - 1s/step - accuracy: 0.7040 - loss: 0.6099 - precision: 0.7569 -
recall: 0.6010 - val_accuracy: 0.6275 - val_loss: 0.6790 - val_precision: 0.5889
- val_recall: 0.8450

Epoch 2/30

13/13 - 18s - 1s/step - accuracy: 0.7375 - loss: 0.5224 - precision: 0.7387 -
recall: 0.7350 - val_accuracy: 0.7150 - val_loss: 0.5928 - val_precision: 0.6807
- val_recall: 0.8100

Epoch 3/30

13/13 - 18s - 1s/step - accuracy: 0.7594 - loss: 0.4923 - precision: 0.7427 -
recall: 0.7937 - val_accuracy: 0.7275 - val_loss: 0.5966 - val_precision: 0.7358
- val_recall: 0.7100

Epoch 4/30

13/13 - 17s - 1s/step - accuracy: 0.8000 - loss: 0.4405 - precision: 0.7850 -
recall: 0.8263 - val_accuracy: 0.7575 - val_loss: 0.5186 - val_precision: 0.7725
- val_recall: 0.7300

Epoch 5/30

13/13 - 18s - 1s/step - accuracy: 0.8050 - loss: 0.4425 - precision: 0.7805 -
recall: 0.8487 - val_accuracy: 0.6600 - val_loss: 0.5973 - val_precision: 0.8333
- val_recall: 0.4000

Epoch 6/30

13/13 - 17s - 1s/step - accuracy: 0.8175 - loss: 0.4151 - precision: 0.7953 -
recall: 0.8550 - val_accuracy: 0.7575 - val_loss: 0.4948 - val_precision: 0.6914
- val_recall: 0.9300

Epoch 7/30
13/13 - 18s - 1s/step - accuracy: 0.8319 - loss: 0.3811 - precision: 0.8105 -
recall: 0.8662 - val_accuracy: 0.7825 - val_loss: 0.4821 - val_precision: 0.7198
- val_recall: 0.9250

Epoch 8/30
13/13 - 18s - 1s/step - accuracy: 0.8431 - loss: 0.3638 - precision: 0.8181 -
recall: 0.8825 - val_accuracy: 0.8125 - val_loss: 0.4028 - val_precision: 0.8019
- val_recall: 0.8300

Epoch 9/30
13/13 - 18s - 1s/step - accuracy: 0.8394 - loss: 0.3594 - precision: 0.8190 -
recall: 0.8712 - val_accuracy: 0.8000 - val_loss: 0.4389 - val_precision: 0.8000
- val_recall: 0.8000

Epoch 10/30
13/13 - 18s - 1s/step - accuracy: 0.8594 - loss: 0.3314 - precision: 0.8410 -
recall: 0.8863 - val_accuracy: 0.8025 - val_loss: 0.4235 - val_precision: 0.8954
- val_recall: 0.6850

Epoch 11/30
13/13 - 18s - 1s/step - accuracy: 0.8444 - loss: 0.3390 - precision: 0.8439 -
recall: 0.8450 - val_accuracy: 0.7925 - val_loss: 0.4465 - val_precision: 0.7773
- val_recall: 0.8200

Epoch 12/30
13/13 - 18s - 1s/step - accuracy: 0.8587 - loss: 0.3273 - precision: 0.8409 -
recall: 0.8850 - val_accuracy: 0.8125 - val_loss: 0.4230 - val_precision: 0.7729
- val_recall: 0.8850

Epoch 13/30
13/13 - 18s - 1s/step - accuracy: 0.8594 - loss: 0.3053 - precision: 0.8419 -
recall: 0.8850 - val_accuracy: 0.7375 - val_loss: 0.5317 - val_precision: 0.8276
- val_recall: 0.6000

Epoch 13: early stopping
Restoring model weights from the end of the best epoch: 8.
evaluating best model from this run

4/4 1s 246ms/step -
accuracy: 0.8078 - loss: 0.3984 - precision: 0.5809 - recall: 0.6519
testing: lr=0.0001, filters_conv1=16, units_dense1=128, dropout_rate=0.4

Epoch 1/30
13/13 - 19s - 1s/step - accuracy: 0.6555 - loss: 0.6638 - precision: 0.6794 -
recall: 0.5890 - val_accuracy: 0.6300 - val_loss: 0.6509 - val_precision: 0.6368
- val_recall: 0.6050

Epoch 2/30
13/13 - 18s - 1s/step - accuracy: 0.6913 - loss: 0.5836 - precision: 0.7068 -
recall: 0.6538 - val_accuracy: 0.6125 - val_loss: 0.6611 - val_precision: 0.6301
- val_recall: 0.5450

Epoch 3/30
13/13 - 18s - 1s/step - accuracy: 0.7169 - loss: 0.5537 - precision: 0.7188 -
recall: 0.7125 - val_accuracy: 0.6250 - val_loss: 0.6601 - val_precision: 0.6157
- val_recall: 0.6650

Epoch 4/30
13/13 - 18s - 1s/step - accuracy: 0.7369 - loss: 0.5473 - precision: 0.7297 -

recall: 0.7525 - val_accuracy: 0.6300 - val_loss: 0.6488 - val_precision: 0.6250
 - val_recall: 0.6500
 Epoch 5/30
 13/13 - 18s - 1s/step - accuracy: 0.7487 - loss: 0.5285 - precision: 0.7403 -
 recall: 0.7663 - val_accuracy: 0.6750 - val_loss: 0.6138 - val_precision: 0.6683
 - val_recall: 0.6950
 Epoch 6/30
 13/13 - 18s - 1s/step - accuracy: 0.7475 - loss: 0.5288 - precision: 0.7409 -
 recall: 0.7613 - val_accuracy: 0.6875 - val_loss: 0.6071 - val_precision: 0.6652
 - val_recall: 0.7550
 Epoch 7/30
 13/13 - 18s - 1s/step - accuracy: 0.7619 - loss: 0.5119 - precision: 0.7422 -
 recall: 0.8025 - val_accuracy: 0.6800 - val_loss: 0.5972 - val_precision: 0.7022
 - val_recall: 0.6250
 Epoch 8/30
 13/13 - 18s - 1s/step - accuracy: 0.7706 - loss: 0.5020 - precision: 0.7599 -
 recall: 0.7912 - val_accuracy: 0.6725 - val_loss: 0.5983 - val_precision: 0.7018
 - val_recall: 0.6000
 Epoch 9/30
 13/13 - 18s - 1s/step - accuracy: 0.7756 - loss: 0.4908 - precision: 0.7603 -
 recall: 0.8050 - val_accuracy: 0.7350 - val_loss: 0.5470 - val_precision: 0.7238
 - val_recall: 0.7600
 Epoch 10/30
 13/13 - 18s - 1s/step - accuracy: 0.7738 - loss: 0.4827 - precision: 0.7724 -
 recall: 0.7763 - val_accuracy: 0.7050 - val_loss: 0.5677 - val_precision: 0.7595
 - val_recall: 0.6000
 Epoch 11/30
 13/13 - 18s - 1s/step - accuracy: 0.7931 - loss: 0.4689 - precision: 0.7795 -
 recall: 0.8175 - val_accuracy: 0.7050 - val_loss: 0.5625 - val_precision: 0.7500
 - val_recall: 0.6150
 Epoch 12/30
 13/13 - 18s - 1s/step - accuracy: 0.7856 - loss: 0.4704 - precision: 0.7698 -
 recall: 0.8150 - val_accuracy: 0.6825 - val_loss: 0.5903 - val_precision: 0.7920
 - val_recall: 0.4950
 Epoch 13/30
 13/13 - 18s - 1s/step - accuracy: 0.7894 - loss: 0.4552 - precision: 0.7689 -
 recall: 0.8275 - val_accuracy: 0.6950 - val_loss: 0.5710 - val_precision: 0.7868
 - val_recall: 0.5350
 Epoch 14/30
 13/13 - 18s - 1s/step - accuracy: 0.8044 - loss: 0.4358 - precision: 0.7944 -
 recall: 0.8213 - val_accuracy: 0.7150 - val_loss: 0.5229 - val_precision: 0.7722
 - val_recall: 0.6100
 Epoch 15/30
 13/13 - 18s - 1s/step - accuracy: 0.7981 - loss: 0.4419 - precision: 0.7891 -
 recall: 0.8138 - val_accuracy: 0.7650 - val_loss: 0.4809 - val_precision: 0.7849
 - val_recall: 0.7300
 Epoch 16/30
 13/13 - 18s - 1s/step - accuracy: 0.8112 - loss: 0.4299 - precision: 0.8007 -

recall: 0.8288 - val_accuracy: 0.7750 - val_loss: 0.4793 - val_precision: 0.8090
- val_recall: 0.7200
Epoch 17/30
13/13 - 18s - 1s/step - accuracy: 0.8087 - loss: 0.4212 - precision: 0.8119 -
recall: 0.8037 - val_accuracy: 0.7825 - val_loss: 0.4552 - val_precision: 0.8054
- val_recall: 0.7450
Epoch 18/30
13/13 - 18s - 1s/step - accuracy: 0.8031 - loss: 0.4284 - precision: 0.7843 -
recall: 0.8363 - val_accuracy: 0.7875 - val_loss: 0.4605 - val_precision: 0.7979
- val_recall: 0.7700
Epoch 19/30
13/13 - 18s - 1s/step - accuracy: 0.8181 - loss: 0.4062 - precision: 0.7991 -
recall: 0.8500 - val_accuracy: 0.7850 - val_loss: 0.4590 - val_precision: 0.8393
- val_recall: 0.7050
Epoch 20/30
13/13 - 18s - 1s/step - accuracy: 0.8131 - loss: 0.3989 - precision: 0.8051 -
recall: 0.8263 - val_accuracy: 0.8075 - val_loss: 0.4514 - val_precision: 0.7709
- val_recall: 0.8750
Epoch 21/30
13/13 - 18s - 1s/step - accuracy: 0.8244 - loss: 0.3950 - precision: 0.8161 -
recall: 0.8375 - val_accuracy: 0.7825 - val_loss: 0.4447 - val_precision: 0.7425
- val_recall: 0.8650
Epoch 22/30
13/13 - 18s - 1s/step - accuracy: 0.8131 - loss: 0.4110 - precision: 0.7923 -
recall: 0.8487 - val_accuracy: 0.7800 - val_loss: 0.4648 - val_precision: 0.7667
- val_recall: 0.8050
Epoch 23/30
13/13 - 18s - 1s/step - accuracy: 0.8294 - loss: 0.3818 - precision: 0.8249 -
recall: 0.8363 - val_accuracy: 0.7625 - val_loss: 0.4765 - val_precision: 0.9268
- val_recall: 0.5700
Epoch 24/30
13/13 - 18s - 1s/step - accuracy: 0.8294 - loss: 0.3761 - precision: 0.8194 -
recall: 0.8450 - val_accuracy: 0.8125 - val_loss: 0.4387 - val_precision: 0.8415
- val_recall: 0.7700
Epoch 25/30
13/13 - 18s - 1s/step - accuracy: 0.8350 - loss: 0.3689 - precision: 0.8198 -
recall: 0.8587 - val_accuracy: 0.8175 - val_loss: 0.4201 - val_precision: 0.8713
- val_recall: 0.7450
Epoch 26/30
13/13 - 18s - 1s/step - accuracy: 0.8250 - loss: 0.3819 - precision: 0.7961 -
recall: 0.8737 - val_accuracy: 0.7800 - val_loss: 0.4500 - val_precision: 0.8415
- val_recall: 0.6900
Epoch 27/30
13/13 - 18s - 1s/step - accuracy: 0.8219 - loss: 0.3824 - precision: 0.8129 -
recall: 0.8363 - val_accuracy: 0.8325 - val_loss: 0.4224 - val_precision: 0.8308
- val_recall: 0.8350
Epoch 28/30
13/13 - 18s - 1s/step - accuracy: 0.8375 - loss: 0.3568 - precision: 0.8140 -

recall: 0.8750 - val_accuracy: 0.8225 - val_loss: 0.4083 - val_precision: 0.7945
 - val_recall: 0.8700
 Epoch 29/30
 13/13 - 18s - 1s/step - accuracy: 0.8275 - loss: 0.3730 - precision: 0.8126 -
 recall: 0.8512 - val_accuracy: 0.8250 - val_loss: 0.3863 - val_precision: 0.7982
 - val_recall: 0.8700
 Epoch 30/30
 13/13 - 18s - 1s/step - accuracy: 0.8275 - loss: 0.3761 - precision: 0.8018 -
 recall: 0.8700 - val_accuracy: 0.8175 - val_loss: 0.4190 - val_precision: 0.8757
 - val_recall: 0.7400
 Restoring model weights from the end of the best epoch: 29.
 evaluating best model from this run
 4/4 1s 251ms/step -
 accuracy: 0.8165 - loss: 0.3983 - precision: 0.5761 - recall: 0.6791
 testing: lr=1e-05, filters_conv1=64, units_dense1=64, dropout_rate=0.4
 Epoch 1/30
 13/13 - 40s - 3s/step - accuracy: 0.5945 - loss: 0.6896 - precision: 0.5990 -
 recall: 0.5720 - val_accuracy: 0.6050 - val_loss: 0.6728 - val_precision: 0.7500
 - val_recall: 0.3150
 Epoch 2/30
 13/13 - 39s - 3s/step - accuracy: 0.6256 - loss: 0.6568 - precision: 0.6476 -
 recall: 0.5512 - val_accuracy: 0.6350 - val_loss: 0.6574 - val_precision: 0.6274
 - val_recall: 0.6650
 Epoch 3/30
 13/13 - 38s - 3s/step - accuracy: 0.6538 - loss: 0.6374 - precision: 0.6420 -
 recall: 0.6950 - val_accuracy: 0.6275 - val_loss: 0.6477 - val_precision: 0.6474
 - val_recall: 0.5600
 Epoch 4/30
 13/13 - 38s - 3s/step - accuracy: 0.6875 - loss: 0.6204 - precision: 0.6923 -
 recall: 0.6750 - val_accuracy: 0.6500 - val_loss: 0.6454 - val_precision: 0.6351
 - val_recall: 0.7050
 Epoch 5/30
 13/13 - 38s - 3s/step - accuracy: 0.6931 - loss: 0.6055 - precision: 0.6939 -
 recall: 0.6913 - val_accuracy: 0.6225 - val_loss: 0.6406 - val_precision: 0.6244
 - val_recall: 0.6150
 Epoch 6/30
 13/13 - 38s - 3s/step - accuracy: 0.6850 - loss: 0.6030 - precision: 0.6775 -
 recall: 0.7063 - val_accuracy: 0.6375 - val_loss: 0.6422 - val_precision: 0.6291
 - val_recall: 0.6700
 Epoch 7/30
 13/13 - 38s - 3s/step - accuracy: 0.7044 - loss: 0.5946 - precision: 0.7051 -
 recall: 0.7025 - val_accuracy: 0.6300 - val_loss: 0.6372 - val_precision: 0.6398
 - val_recall: 0.5950
 Epoch 8/30
 13/13 - 38s - 3s/step - accuracy: 0.7131 - loss: 0.5894 - precision: 0.7183 -
 recall: 0.7013 - val_accuracy: 0.6500 - val_loss: 0.6341 - val_precision: 0.6429
 - val_recall: 0.6750
 Epoch 9/30

13/13 - 38s - 3s/step - accuracy: 0.6981 - loss: 0.5828 - precision: 0.6964 -
recall: 0.7025 - val_accuracy: 0.6350 - val_loss: 0.6321 - val_precision: 0.6552
- val_recall: 0.5700
Epoch 10/30
13/13 - 38s - 3s/step - accuracy: 0.7156 - loss: 0.5767 - precision: 0.7061 -
recall: 0.7387 - val_accuracy: 0.6475 - val_loss: 0.6252 - val_precision: 0.6595
- val_recall: 0.6100
Epoch 11/30
13/13 - 38s - 3s/step - accuracy: 0.7200 - loss: 0.5685 - precision: 0.7245 -
recall: 0.7100 - val_accuracy: 0.6450 - val_loss: 0.6243 - val_precision: 0.6629
- val_recall: 0.5900
Epoch 12/30
13/13 - 38s - 3s/step - accuracy: 0.7025 - loss: 0.5714 - precision: 0.6966 -
recall: 0.7175 - val_accuracy: 0.6450 - val_loss: 0.6196 - val_precision: 0.6629
- val_recall: 0.5900
Epoch 13/30
13/13 - 38s - 3s/step - accuracy: 0.7219 - loss: 0.5593 - precision: 0.7216 -
recall: 0.7225 - val_accuracy: 0.6450 - val_loss: 0.6167 - val_precision: 0.6686
- val_recall: 0.5750
Epoch 14/30
13/13 - 38s - 3s/step - accuracy: 0.7244 - loss: 0.5706 - precision: 0.7155 -
recall: 0.7450 - val_accuracy: 0.6525 - val_loss: 0.6136 - val_precision: 0.6743
- val_recall: 0.5900
Epoch 15/30
13/13 - 38s - 3s/step - accuracy: 0.7331 - loss: 0.5477 - precision: 0.7266 -
recall: 0.7475 - val_accuracy: 0.6450 - val_loss: 0.6196 - val_precision: 0.6859
- val_recall: 0.5350
Epoch 16/30
13/13 - 38s - 3s/step - accuracy: 0.7337 - loss: 0.5512 - precision: 0.7149 -
recall: 0.7775 - val_accuracy: 0.6400 - val_loss: 0.6172 - val_precision: 0.6591
- val_recall: 0.5800
Epoch 17/30
13/13 - 38s - 3s/step - accuracy: 0.7300 - loss: 0.5585 - precision: 0.7323 -
recall: 0.7250 - val_accuracy: 0.6400 - val_loss: 0.6202 - val_precision: 0.6750
- val_recall: 0.5400
Epoch 18/30
13/13 - 38s - 3s/step - accuracy: 0.7437 - loss: 0.5331 - precision: 0.7273 -
recall: 0.7800 - val_accuracy: 0.6600 - val_loss: 0.6107 - val_precision: 0.6667
- val_recall: 0.6400
Epoch 19/30
13/13 - 39s - 3s/step - accuracy: 0.7394 - loss: 0.5414 - precision: 0.7235 -
recall: 0.7750 - val_accuracy: 0.6700 - val_loss: 0.6031 - val_precision: 0.6889
- val_recall: 0.6200
Epoch 20/30
13/13 - 37s - 3s/step - accuracy: 0.7306 - loss: 0.5537 - precision: 0.7220 -
recall: 0.7500 - val_accuracy: 0.6425 - val_loss: 0.6039 - val_precision: 0.6770
- val_recall: 0.5450
Epoch 21/30

13/13 - 37s - 3s/step - accuracy: 0.7513 - loss: 0.5304 - precision: 0.7399 -
recall: 0.7750 - val_accuracy: 0.6625 - val_loss: 0.5943 - val_precision: 0.6836
- val_recall: 0.6050

Epoch 22/30

13/13 - 38s - 3s/step - accuracy: 0.7419 - loss: 0.5404 - precision: 0.7201 -
recall: 0.7912 - val_accuracy: 0.6400 - val_loss: 0.6025 - val_precision: 0.6772
- val_recall: 0.5350

Epoch 23/30

13/13 - 38s - 3s/step - accuracy: 0.7450 - loss: 0.5329 - precision: 0.7300 -
recall: 0.7775 - val_accuracy: 0.6600 - val_loss: 0.5972 - val_precision: 0.6839
- val_recall: 0.5950

Epoch 24/30

13/13 - 38s - 3s/step - accuracy: 0.7444 - loss: 0.5372 - precision: 0.7234 -
recall: 0.7912 - val_accuracy: 0.6500 - val_loss: 0.6000 - val_precision: 0.6852
- val_recall: 0.5550

Epoch 25/30

13/13 - 37s - 3s/step - accuracy: 0.7369 - loss: 0.5355 - precision: 0.7181 -
recall: 0.7800 - val_accuracy: 0.6400 - val_loss: 0.6004 - val_precision: 0.6842
- val_recall: 0.5200

Epoch 26/30

13/13 - 37s - 3s/step - accuracy: 0.7494 - loss: 0.5273 - precision: 0.7418 -
recall: 0.7650 - val_accuracy: 0.6650 - val_loss: 0.5865 - val_precision: 0.6875
- val_recall: 0.6050

Epoch 27/30

13/13 - 38s - 3s/step - accuracy: 0.7531 - loss: 0.5205 - precision: 0.7273 -
recall: 0.8100 - val_accuracy: 0.6500 - val_loss: 0.6006 - val_precision: 0.7027
- val_recall: 0.5200

Epoch 28/30

13/13 - 37s - 3s/step - accuracy: 0.7444 - loss: 0.5346 - precision: 0.7314 -
recall: 0.7725 - val_accuracy: 0.6575 - val_loss: 0.5960 - val_precision: 0.6864
- val_recall: 0.5800

Epoch 29/30

13/13 - 38s - 3s/step - accuracy: 0.7481 - loss: 0.5189 - precision: 0.7300 -
recall: 0.7875 - val_accuracy: 0.6675 - val_loss: 0.5848 - val_precision: 0.6914
- val_recall: 0.6050

Epoch 30/30

13/13 - 37s - 3s/step - accuracy: 0.7538 - loss: 0.5086 - precision: 0.7494 -
recall: 0.7625 - val_accuracy: 0.6425 - val_loss: 0.5860 - val_precision: 0.6727
- val_recall: 0.5550

Restoring model weights from the end of the best epoch: 29.

evaluating best model from this run

4/4 2s 492ms/step -

accuracy: 0.6938 - loss: 0.5609 - precision: 0.4950 - recall: 0.5019

testing: lr=0.001, filters_conv1=32, units_dense1=128, dropout_rate=0.2

Epoch 1/30

13/13 - 24s - 2s/step - accuracy: 0.5750 - loss: 1.0871 - precision: 0.6280 -
recall: 0.3680 - val_accuracy: 0.6225 - val_loss: 0.6555 - val_precision: 0.6269
- val_recall: 0.6050

Epoch 2/30

13/13 - 23s - 2s/step - accuracy: 0.7169 - loss: 0.5723 - precision: 0.7188 -
recall: 0.7125 - val_accuracy: 0.6400 - val_loss: 0.6373 - val_precision: 0.6818
- val_recall: 0.5250

Epoch 3/30

13/13 - 23s - 2s/step - accuracy: 0.7425 - loss: 0.5161 - precision: 0.7407 -
recall: 0.7462 - val_accuracy: 0.6775 - val_loss: 0.6110 - val_precision: 0.6437
- val_recall: 0.7950

Epoch 4/30

13/13 - 23s - 2s/step - accuracy: 0.7700 - loss: 0.4887 - precision: 0.7541 -
recall: 0.8012 - val_accuracy: 0.7025 - val_loss: 0.5842 - val_precision: 0.7263
- val_recall: 0.6500

Epoch 5/30

13/13 - 23s - 2s/step - accuracy: 0.7831 - loss: 0.4754 - precision: 0.7542 -
recall: 0.8400 - val_accuracy: 0.7525 - val_loss: 0.5188 - val_precision: 0.7538
- val_recall: 0.7500

Epoch 6/30

13/13 - 23s - 2s/step - accuracy: 0.7975 - loss: 0.4373 - precision: 0.7874 -
recall: 0.8150 - val_accuracy: 0.7850 - val_loss: 0.4912 - val_precision: 0.7500
- val_recall: 0.8550

Epoch 7/30

13/13 - 23s - 2s/step - accuracy: 0.7956 - loss: 0.4415 - precision: 0.7799 -
recall: 0.8238 - val_accuracy: 0.8025 - val_loss: 0.5136 - val_precision: 0.7410
- val_recall: 0.9300

Epoch 8/30

13/13 - 23s - 2s/step - accuracy: 0.8144 - loss: 0.4157 - precision: 0.7976 -
recall: 0.8425 - val_accuracy: 0.7000 - val_loss: 0.5387 - val_precision: 0.8846
- val_recall: 0.4600

Epoch 9/30

13/13 - 24s - 2s/step - accuracy: 0.8263 - loss: 0.3957 - precision: 0.8115 -
recall: 0.8500 - val_accuracy: 0.8100 - val_loss: 0.4560 - val_precision: 0.7627
- val_recall: 0.9000

Epoch 10/30

13/13 - 24s - 2s/step - accuracy: 0.8375 - loss: 0.3783 - precision: 0.8075 -
recall: 0.8863 - val_accuracy: 0.7975 - val_loss: 0.4482 - val_precision: 0.7990
- val_recall: 0.7950

Epoch 11/30

13/13 - 23s - 2s/step - accuracy: 0.8338 - loss: 0.3729 - precision: 0.8240 -
recall: 0.8487 - val_accuracy: 0.7875 - val_loss: 0.4441 - val_precision: 0.8402
- val_recall: 0.7100

Epoch 12/30

13/13 - 23s - 2s/step - accuracy: 0.8375 - loss: 0.3565 - precision: 0.8342 -
recall: 0.8425 - val_accuracy: 0.8300 - val_loss: 0.4056 - val_precision: 0.8267
- val_recall: 0.8350

Epoch 13/30

13/13 - 23s - 2s/step - accuracy: 0.8644 - loss: 0.3354 - precision: 0.8474 -
recall: 0.8888 - val_accuracy: 0.8425 - val_loss: 0.4078 - val_precision: 0.8703
- val_recall: 0.8050

Epoch 14/30

13/13 - 23s - 2s/step - accuracy: 0.8525 - loss: 0.3345 - precision: 0.8341 -
recall: 0.8800 - val_accuracy: 0.8275 - val_loss: 0.4005 - val_precision: 0.8701
- val_recall: 0.7700

Epoch 15/30

13/13 - 23s - 2s/step - accuracy: 0.8550 - loss: 0.3139 - precision: 0.8541 -
recall: 0.8562 - val_accuracy: 0.8275 - val_loss: 0.3906 - val_precision: 0.7991
- val_recall: 0.8750

Epoch 16/30

13/13 - 23s - 2s/step - accuracy: 0.8569 - loss: 0.3379 - precision: 0.8371 -
recall: 0.8863 - val_accuracy: 0.7025 - val_loss: 0.5101 - val_precision: 0.8857
- val_recall: 0.4650

Epoch 17/30

13/13 - 23s - 2s/step - accuracy: 0.8594 - loss: 0.3152 - precision: 0.8394 -
recall: 0.8888 - val_accuracy: 0.8225 - val_loss: 0.4209 - val_precision: 0.8862
- val_recall: 0.7400

Epoch 18/30

13/13 - 23s - 2s/step - accuracy: 0.8706 - loss: 0.2952 - precision: 0.8585 -
recall: 0.8875 - val_accuracy: 0.8625 - val_loss: 0.3538 - val_precision: 0.8836
- val_recall: 0.8350

Epoch 19/30

13/13 - 23s - 2s/step - accuracy: 0.8587 - loss: 0.3134 - precision: 0.8475 -
recall: 0.8750 - val_accuracy: 0.8575 - val_loss: 0.3632 - val_precision: 0.8667
- val_recall: 0.8450

Epoch 20/30

13/13 - 23s - 2s/step - accuracy: 0.8700 - loss: 0.2817 - precision: 0.8619 -
recall: 0.8813 - val_accuracy: 0.8375 - val_loss: 0.3865 - val_precision: 0.8814
- val_recall: 0.7800

Epoch 21/30

13/13 - 23s - 2s/step - accuracy: 0.8769 - loss: 0.2865 - precision: 0.8585 -
recall: 0.9025 - val_accuracy: 0.8500 - val_loss: 0.3639 - val_precision: 0.8431
- val_recall: 0.8600

Epoch 22/30

13/13 - 23s - 2s/step - accuracy: 0.8875 - loss: 0.2617 - precision: 0.8780 -
recall: 0.9000 - val_accuracy: 0.8600 - val_loss: 0.3447 - val_precision: 0.8600
- val_recall: 0.8600

Epoch 23/30

13/13 - 23s - 2s/step - accuracy: 0.8712 - loss: 0.2950 - precision: 0.8578 -
recall: 0.8900 - val_accuracy: 0.8675 - val_loss: 0.3319 - val_precision: 0.8517
- val_recall: 0.8900

Epoch 24/30

13/13 - 23s - 2s/step - accuracy: 0.8775 - loss: 0.2848 - precision: 0.8595 -
recall: 0.9025 - val_accuracy: 0.8650 - val_loss: 0.3723 - val_precision: 0.8883
- val_recall: 0.8350

Epoch 25/30

13/13 - 23s - 2s/step - accuracy: 0.8913 - loss: 0.2547 - precision: 0.8826 -
recall: 0.9025 - val_accuracy: 0.8850 - val_loss: 0.3166 - val_precision: 0.9010
- val_recall: 0.8650

Epoch 26/30
13/13 - 23s - 2s/step - accuracy: 0.9094 - loss: 0.2371 - precision: 0.8970 -
recall: 0.9250 - val_accuracy: 0.8700 - val_loss: 0.3426 - val_precision: 0.9205
- val_recall: 0.8100

Epoch 27/30
13/13 - 23s - 2s/step - accuracy: 0.8881 - loss: 0.2535 - precision: 0.8819 -
recall: 0.8963 - val_accuracy: 0.8750 - val_loss: 0.3154 - val_precision: 0.9213
- val_recall: 0.8200

Epoch 28/30
13/13 - 23s - 2s/step - accuracy: 0.8844 - loss: 0.2605 - precision: 0.8773 -
recall: 0.8938 - val_accuracy: 0.8600 - val_loss: 0.3395 - val_precision: 0.9235
- val_recall: 0.7850

Epoch 29/30
13/13 - 24s - 2s/step - accuracy: 0.8888 - loss: 0.2636 - precision: 0.8859 -
recall: 0.8925 - val_accuracy: 0.8700 - val_loss: 0.3254 - val_precision: 0.8895
- val_recall: 0.8450

Epoch 30/30
13/13 - 23s - 2s/step - accuracy: 0.8906 - loss: 0.2484 - precision: 0.8834 -
recall: 0.9000 - val_accuracy: 0.8825 - val_loss: 0.2855 - val_precision: 0.9005
- val_recall: 0.8600

Restoring model weights from the end of the best epoch: 30.
evaluating best model from this run

4/4 2s 338ms/step -
accuracy: 0.8871 - loss: 0.2705 - precision: 0.6804 - recall: 0.6811
testing: lr=1e-05, filters_conv1=16, units_dense1=256, dropout_rate=0.3

Epoch 1/30
13/13 - 20s - 2s/step - accuracy: 0.7115 - loss: 0.6194 - precision: 0.7311 -
recall: 0.6690 - val_accuracy: 0.6200 - val_loss: 0.6825 - val_precision: 0.6053
- val_recall: 0.6900

Epoch 2/30
13/13 - 18s - 1s/step - accuracy: 0.7381 - loss: 0.5477 - precision: 0.7408 -
recall: 0.7325 - val_accuracy: 0.6050 - val_loss: 0.6734 - val_precision: 0.6129
- val_recall: 0.5700

Epoch 3/30
13/13 - 18s - 1s/step - accuracy: 0.7212 - loss: 0.5481 - precision: 0.7287 -
recall: 0.7050 - val_accuracy: 0.6200 - val_loss: 0.6405 - val_precision: 0.6446
- val_recall: 0.5350

Epoch 4/30
13/13 - 18s - 1s/step - accuracy: 0.7631 - loss: 0.5091 - precision: 0.7615 -
recall: 0.7663 - val_accuracy: 0.6450 - val_loss: 0.6259 - val_precision: 0.6686
- val_recall: 0.5750

Epoch 5/30
13/13 - 18s - 1s/step - accuracy: 0.7625 - loss: 0.5044 - precision: 0.7672 -
recall: 0.7538 - val_accuracy: 0.6250 - val_loss: 0.6304 - val_precision: 0.6667
- val_recall: 0.5000

Epoch 6/30
13/13 - 18s - 1s/step - accuracy: 0.7581 - loss: 0.5139 - precision: 0.7604 -
recall: 0.7538 - val_accuracy: 0.6625 - val_loss: 0.6049 - val_precision: 0.6816

- val_recall: 0.6100
Epoch 7/30
13/13 - 18s - 1s/step - accuracy: 0.7669 - loss: 0.4867 - precision: 0.7515 -
recall: 0.7975 - val_accuracy: 0.6675 - val_loss: 0.5859 - val_precision: 0.7055
- val_recall: 0.5750
Epoch 8/30
13/13 - 18s - 1s/step - accuracy: 0.7769 - loss: 0.4920 - precision: 0.7646 -
recall: 0.8000 - val_accuracy: 0.6875 - val_loss: 0.5775 - val_precision: 0.7551
- val_recall: 0.5550
Epoch 9/30
13/13 - 18s - 1s/step - accuracy: 0.7675 - loss: 0.4895 - precision: 0.7716 -
recall: 0.7600 - val_accuracy: 0.7300 - val_loss: 0.5494 - val_precision: 0.7614
- val_recall: 0.6700
Epoch 10/30
13/13 - 18s - 1s/step - accuracy: 0.7769 - loss: 0.4775 - precision: 0.7597 -
recall: 0.8100 - val_accuracy: 0.7300 - val_loss: 0.5456 - val_precision: 0.7614
- val_recall: 0.6700
Epoch 11/30
13/13 - 18s - 1s/step - accuracy: 0.7844 - loss: 0.4676 - precision: 0.7771 -
recall: 0.7975 - val_accuracy: 0.7275 - val_loss: 0.5388 - val_precision: 0.7758
- val_recall: 0.6400
Epoch 12/30
13/13 - 18s - 1s/step - accuracy: 0.7962 - loss: 0.4533 - precision: 0.7985 -
recall: 0.7925 - val_accuracy: 0.7000 - val_loss: 0.5457 - val_precision: 0.7941
- val_recall: 0.5400
Epoch 13/30
13/13 - 18s - 1s/step - accuracy: 0.7744 - loss: 0.4698 - precision: 0.7654 -
recall: 0.7912 - val_accuracy: 0.6775 - val_loss: 0.5693 - val_precision: 0.7934
- val_recall: 0.4800
Epoch 14/30
13/13 - 18s - 1s/step - accuracy: 0.8012 - loss: 0.4560 - precision: 0.7783 -
recall: 0.8425 - val_accuracy: 0.6675 - val_loss: 0.6049 - val_precision: 0.8941
- val_recall: 0.3800
Epoch 15/30
13/13 - 18s - 1s/step - accuracy: 0.7906 - loss: 0.4614 - precision: 0.7888 -
recall: 0.7937 - val_accuracy: 0.7425 - val_loss: 0.5306 - val_precision: 0.7594
- val_recall: 0.7100
Epoch 16/30
13/13 - 18s - 1s/step - accuracy: 0.7950 - loss: 0.4453 - precision: 0.7843 -
recall: 0.8138 - val_accuracy: 0.7625 - val_loss: 0.5203 - val_precision: 0.7512
- val_recall: 0.7850
Epoch 17/30
13/13 - 18s - 1s/step - accuracy: 0.8106 - loss: 0.4408 - precision: 0.8027 -
recall: 0.8238 - val_accuracy: 0.7575 - val_loss: 0.5185 - val_precision: 0.7668
- val_recall: 0.7400
Epoch 18/30
13/13 - 18s - 1s/step - accuracy: 0.8081 - loss: 0.4325 - precision: 0.7938 -
recall: 0.8325 - val_accuracy: 0.7600 - val_loss: 0.5000 - val_precision: 0.7653

- val_recall: 0.7500

Epoch 19/30

13/13 - 18s - 1s/step - accuracy: 0.8006 - loss: 0.4410 - precision: 0.7929 - recall: 0.8138 - val_accuracy: 0.7575 - val_loss: 0.5006 - val_precision: 0.8601 - val_recall: 0.6150

Epoch 20/30

13/13 - 18s - 1s/step - accuracy: 0.8125 - loss: 0.4181 - precision: 0.8109 - recall: 0.8150 - val_accuracy: 0.7650 - val_loss: 0.4985 - val_precision: 0.8193 - val_recall: 0.6800

Epoch 21/30

13/13 - 18s - 1s/step - accuracy: 0.8125 - loss: 0.4332 - precision: 0.8086 - recall: 0.8188 - val_accuracy: 0.7825 - val_loss: 0.4956 - val_precision: 0.8266 - val_recall: 0.7150

Epoch 22/30

13/13 - 18s - 1s/step - accuracy: 0.8131 - loss: 0.4204 - precision: 0.7944 - recall: 0.8450 - val_accuracy: 0.7150 - val_loss: 0.5418 - val_precision: 0.8772 - val_recall: 0.5000

Epoch 23/30

13/13 - 18s - 1s/step - accuracy: 0.8075 - loss: 0.4351 - precision: 0.7874 - recall: 0.8425 - val_accuracy: 0.7700 - val_loss: 0.4965 - val_precision: 0.8462 - val_recall: 0.6600

Epoch 24/30

13/13 - 18s - 1s/step - accuracy: 0.8087 - loss: 0.4152 - precision: 0.8159 - recall: 0.7975 - val_accuracy: 0.7750 - val_loss: 0.4853 - val_precision: 0.7723 - val_recall: 0.7800

Epoch 25/30

13/13 - 18s - 1s/step - accuracy: 0.8125 - loss: 0.4127 - precision: 0.7990 - recall: 0.8350 - val_accuracy: 0.7825 - val_loss: 0.4923 - val_precision: 0.8792 - val_recall: 0.6550

Epoch 26/30

13/13 - 18s - 1s/step - accuracy: 0.8150 - loss: 0.4060 - precision: 0.7986 - recall: 0.8425 - val_accuracy: 0.7875 - val_loss: 0.4814 - val_precision: 0.8859 - val_recall: 0.6600

Epoch 27/30

13/13 - 18s - 1s/step - accuracy: 0.8275 - loss: 0.3966 - precision: 0.8164 - recall: 0.8450 - val_accuracy: 0.7700 - val_loss: 0.4867 - val_precision: 0.8699 - val_recall: 0.6350

Epoch 28/30

13/13 - 18s - 1s/step - accuracy: 0.8281 - loss: 0.4007 - precision: 0.8285 - recall: 0.8275 - val_accuracy: 0.7500 - val_loss: 0.5064 - val_precision: 0.8968 - val_recall: 0.5650

Epoch 29/30

13/13 - 18s - 1s/step - accuracy: 0.8188 - loss: 0.4003 - precision: 0.8087 - recall: 0.8350 - val_accuracy: 0.7650 - val_loss: 0.4743 - val_precision: 0.7849 - val_recall: 0.7300

Epoch 30/30

13/13 - 18s - 1s/step - accuracy: 0.8175 - loss: 0.4019 - precision: 0.8151 - recall: 0.8213 - val_accuracy: 0.7975 - val_loss: 0.4629 - val_precision: 0.8362

```
- val_recall: 0.7400
Restoring model weights from the end of the best epoch: 30.
evaluating best model from this run
4/4          1s 265ms/step -
accuracy: 0.8172 - loss: 0.4547 - precision: 0.6153 - recall: 0.5867
hyperparameter tuning finished.
```

```
[21]: # convert results to dataframe
results_df = pd.DataFrame(results_list)

# sort by validation accuracy (descending)
results_df = results_df.sort_values(by="val_recall", ascending=False)

# print results
print("\ntuning results summary:")
print(results_df)

# save results to csv
results_filename = "../results/hyperparameter_tuning_results_random_search.csv"
results_df.to_csv(results_filename, index=False)
print(f"\nresults saved to {results_filename}")
```

tuning results summary:

	learning_rate	filters_conv1	units_dense1	dropout_rate	val_loss \
1	0.00010	16	128	0.4	0.386285
3	0.00100	32	128	0.2	0.285464
0	0.00100	16	64	0.2	0.402848
4	0.00001	16	256	0.3	0.462921
2	0.00001	64	64	0.4	0.584815

	val_accuracy	val_precision	val_recall	epochs_trained
1	0.8250	0.798165	0.870	30
3	0.8825	0.900524	0.860	30
0	0.8125	0.801932	0.830	13
4	0.7975	0.836158	0.740	30
2	0.6675	0.691429	0.605	30

results saved to ../results/hyperparameter_tuning_results_random_search.csv

```
[22]: # concatenate train and val data
X_all = []
y_all = []

train_data_gen.reset()
val_data_gen.reset()

for batch_x, batch_y in train_data_gen:
```

```

X_all.append(batch_x)
y_all.append(batch_y)
if len(X_all) * BATCH_SIZE >= train_data_gen.samples:
    break

for batch_x, batch_y in val_data_gen:
    X_all.append(batch_x)
    y_all.append(batch_y)
    if len(X_all) * BATCH_SIZE >= val_data_gen.samples + train_data_gen.samples:
        break

X_all = np.concatenate(X_all)
y_all = np.concatenate(y_all)

# create dataset
train_val_ds = (
    tf.data.Dataset.from_tensor_slices((X_all, y_all))
    .shuffle(1000)
    .batch(BATCH_SIZE)
    .prefetch(tf.data.AUTOTUNE)
)

```

```

[23]: print("\nretraining the best model on the combined train+validation data")
best_params = results_df.iloc[0]
best_lr = best_params["learning_rate"]
best_filters_conv1 = int(
    best_params["filters_conv1"]
) # cast to int, wasn't working before
best_units_dense1 = int(
    best_params["units_dense1"]
) # cast to int, wasn't working before
best_dropout_rate = best_params["dropout_rate"]
print(
    f"best hyperparameters found: lr={best_lr},
    ↳filters_conv1={best_filters_conv1}, units_dense1={best_units_dense1},
    ↳dropout_rate={best_dropout_rate}"
)

```

retraining the best model on the combined train+validation data
best hyperparameters found: lr=0.0001, filters_conv1=16, units_dense1=128,
dropout_rate=0.4

```

[24]: # build the best model
final_model = build_model(
    input_shape,
    filters_conv1=best_filters_conv1,

```



```

        units_dense1=best_units_dense1,
        dropout_rate=best_dropout_rate,
    )

    # compile the best model
    optimizer = tf.keras.optimizers.Adam(learning_rate=best_lr)
    final_model.compile(optimizer=optimizer, loss="binary_crossentropy",
        metrics=metrics)

    # define steps per epoch
    steps_per_epoch = (train_data_gen.samples + val_data_gen.samples) // BATCH_SIZE

    # train the best model on the combined train+val data
    final_history = final_model.fit(
        train_val_ds,
        epochs=EPOCHS,
        steps_per_epoch=steps_per_epoch,
        verbose=2,
    )

```

Epoch 1/30

15/15 - 20s - 1s/step - accuracy: 0.6328 - loss: 0.6531 - precision: 0.6823 - recall: 0.4826

Epoch 2/30

15/15 - 1s - 40ms/step - accuracy: 0.6719 - loss: 0.6383 - precision: 0.7143 - recall: 0.6944

Epoch 3/30

/opt/anaconda3/envs/ml-2025/lib/python3.12/site-packages/keras/src/trainers/epoch_iterator.py:107: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()` function when building your dataset.

self._interrupted_warning()

15/15 - 18s - 1s/step - accuracy: 0.6958 - loss: 0.6035 - precision: 0.7222 - recall: 0.6314

Epoch 4/30

15/15 - 1s - 41ms/step - accuracy: 0.7031 - loss: 0.6559 - precision: 0.6562 - recall: 0.7241

Epoch 5/30

15/15 - 18s - 1s/step - accuracy: 0.7130 - loss: 0.5663 - precision: 0.7313 - recall: 0.6677

Epoch 6/30

15/15 - 1s - 41ms/step - accuracy: 0.7656 - loss: 0.5726 - precision: 0.7778 - recall: 0.7000

Epoch 7/30

15/15 - 18s - 1s/step - accuracy: 0.7260 - loss: 0.5446 - precision: 0.7665 - recall: 0.6461

Epoch 8/30
15/15 - 1s - 40ms/step - accuracy: 0.7656 - loss: 0.4862 - precision: 0.6944 - recall: 0.8621
Epoch 9/30
15/15 - 18s - 1s/step - accuracy: 0.7391 - loss: 0.5077 - precision: 0.7494 - recall: 0.7125
Epoch 10/30
15/15 - 1s - 41ms/step - accuracy: 0.7344 - loss: 0.5241 - precision: 0.6842 - recall: 0.8387
Epoch 11/30
15/15 - 18s - 1s/step - accuracy: 0.7573 - loss: 0.4960 - precision: 0.7702 - recall: 0.7247
Epoch 12/30
15/15 - 1s - 41ms/step - accuracy: 0.7188 - loss: 0.5379 - precision: 0.7647 - recall: 0.7222
Epoch 13/30
15/15 - 18s - 1s/step - accuracy: 0.7896 - loss: 0.4661 - precision: 0.7910 - recall: 0.7811
Epoch 14/30
15/15 - 1s - 41ms/step - accuracy: 0.8281 - loss: 0.4705 - precision: 0.8485 - recall: 0.8235
Epoch 15/30
15/15 - 18s - 1s/step - accuracy: 0.8068 - loss: 0.4287 - precision: 0.8229 - recall: 0.7751
Epoch 16/30
15/15 - 1s - 40ms/step - accuracy: 0.7812 - loss: 0.4602 - precision: 0.7949 - recall: 0.8378
Epoch 17/30
15/15 - 18s - 1s/step - accuracy: 0.8125 - loss: 0.4210 - precision: 0.8083 - recall: 0.8168
Epoch 18/30
15/15 - 1s - 42ms/step - accuracy: 0.7812 - loss: 0.3636 - precision: 0.8261 - recall: 0.6552
Epoch 19/30
15/15 - 18s - 1s/step - accuracy: 0.8302 - loss: 0.3909 - precision: 0.8256 - recall: 0.8309
Epoch 20/30
15/15 - 1s - 40ms/step - accuracy: 0.9062 - loss: 0.3380 - precision: 0.9211 - recall: 0.9211
Epoch 21/30
15/15 - 18s - 1s/step - accuracy: 0.8458 - loss: 0.3711 - precision: 0.8370 - recall: 0.8529
Epoch 22/30
15/15 - 1s - 40ms/step - accuracy: 0.8906 - loss: 0.2683 - precision: 0.9706 - recall: 0.8462
Epoch 23/30
15/15 - 18s - 1s/step - accuracy: 0.8620 - loss: 0.3395 - precision: 0.8591 - recall: 0.8636

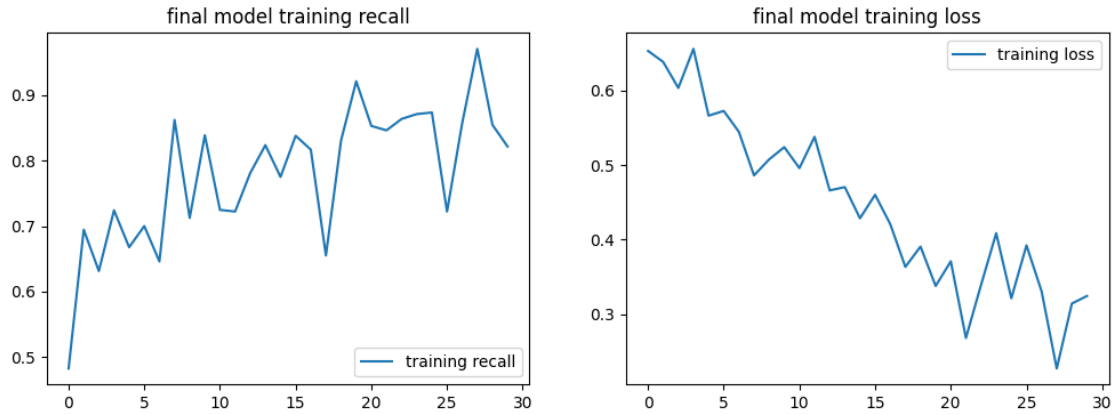
Epoch 24/30
15/15 - 1s - 41ms/step - accuracy: 0.8594 - loss: 0.4087 - precision: 0.8438 - recall: 0.8710
Epoch 25/30
15/15 - 18s - 1s/step - accuracy: 0.8698 - loss: 0.3215 - precision: 0.8643 - recall: 0.8734
Epoch 26/30
15/15 - 1s - 42ms/step - accuracy: 0.7812 - loss: 0.3924 - precision: 0.8667 - recall: 0.7222
Epoch 27/30
15/15 - 19s - 1s/step - accuracy: 0.8630 - loss: 0.3308 - precision: 0.8650 - recall: 0.8568
Epoch 28/30
15/15 - 1s - 41ms/step - accuracy: 0.9375 - loss: 0.2274 - precision: 0.9167 - recall: 0.9706
Epoch 29/30
15/15 - 18s - 1s/step - accuracy: 0.8719 - loss: 0.3145 - precision: 0.8842 - recall: 0.8546
Epoch 30/30
15/15 - 1s - 41ms/step - accuracy: 0.8438 - loss: 0.3246 - precision: 0.8214 - recall: 0.8214

```
[25]: print("\nplotting final training history")
      final_recall = final_history.history["recall"]
      final_loss = final_history.history["loss"]

      plt.figure(figsize=(12, 4))
      plt.subplot(1, 2, 1)
      plt.plot(epochs_range, final_recall, label="training recall")
      plt.legend(loc="lower right")
      plt.title("final model training recall")

      plt.subplot(1, 2, 2)
      plt.plot(epochs_range, final_loss, label="training loss")
      plt.legend(loc="upper right")
      plt.title("final model training loss")
      plt.show()
```

plotting final training history



```
[26]: model_save_path = "../models/baseline_model.keras"
print(f"\nsaving final model to {model_save_path}")
final_model.save(model_save_path)

print("loading model back")
loaded_model = tf.keras.models.load_model(model_save_path)
print("model loaded successfully")
```

```
saving final model to ../models/baseline_model.keras
loading model back
model loaded successfully
```

```
[27]: print("\nevaluating loaded model on test data")
test_results = loaded_model.evaluate(test_data_gen, verbose=1)
print(f"test loss: {test_results[0]}")
print(f"test accuracy: {test_results[1]}")
print(f"test precision: {test_results[2]}")
print(f"test recall: {test_results[3]}")
```

```
evaluating loaded model on test data
```

```
/opt/anaconda3/envs/ml-2025/lib/python3.12/site-
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
```


```
self._warn_if_super_not_called()
```

```
2/2          1s 220ms/step -
accuracy: 0.8346 - loss: 0.3719 - precision: 0.7460 - recall: 0.8695
test loss: 0.372536838054657
```

test accuracy: 0.8299999833106995
test precision: 0.8235294222831726
test recall: 0.8399999737739563

```
[28]: print("\ngenerating confusion matrix")
      # get predictions (probabilities)
      y_pred_prob = loaded_model.predict(test_data_gen)
      # convert probabilities to binary predictions
      y_pred = (y_pred_prob > 0.5).astype(int).flatten()
      # get true labels
      y_true = test_data_gen.classes
      # calculate confusion matrix
      cm = confusion_matrix(y_true, y_pred)
      # plot confusion matrix
      plt.figure(figsize=(6, 5))
      sns.heatmap(
          cm,
          annot=True,
          fmt="d",
          cmap="Blues",
          xticklabels=class_names,
          yticklabels=class_names,
      )
      plt.xlabel("predicted label")
      plt.ylabel("true label")
      plt.title("confusion matrix on test set")
      plt.show()
```

generating confusion matrix
2/2 1s 229ms/step



```
[29]: print("\nplotting sample test images with predictions")
# get a batch of raw (unnormalized) images and labels
images_raw, labels_raw = next(iter(test_data_gen_raw))
# get predictions for this batch using the normalized data generator
# important: ensure we use the *same batch* for predictions
# we need to reset the generator to be sure we get the same batch
test_data_gen.reset()
images_norm, _ = next(iter(test_data_gen)) # get normalized images for
↳ prediction
batch_pred_prob = loaded_model.predict(images_norm)
batch_pred = (batch_pred_prob > 0.5).astype(int).flatten()

plt.figure(figsize=(10, 10))
for i in range(9): # display 9 samples
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(images_raw[i].astype("uint8")) # display raw image
    true_label = class_names[int(labels_raw[i])] # cast to int
    pred_label = class_names[batch_pred[i]]
    prob = batch_pred_prob[i][0]
```

```
plt.title(f"true: {true_label}\npred: {pred_label} ({prob:.2f})")
plt.axis("off")
plt.tight_layout()
plt.show()
```

plotting sample test images with predictions
4/4 0s 88ms/step

