

## resnet

April 22, 2025

```
[1]: import os
import itertools
import random
import numpy as np
import tensorflow as tf
import pandas as pd
import matplotlib.pyplot as plt
from resnet import (
    build_transfer_model,
    train_model,
    fine_tune_model,
    plot_training_history,
    plot_confusion_matrix,
    PROCESSED_DIR,
    IMG_HEIGHT,
    IMG_WIDTH,
    NUM_CHANNELS,
)
import data_pipeline as pipeline
```

```
[2]: BATCH_SIZE = 32  # smaller batch size for transfer learning
SEED = 42
np.random.seed(SEED)
tf.random.set_seed(SEED)
```

```
[3]: # get image paths, mean and std
train_dir = os.path.join(PROCESSED_DIR, "train")
val_dir = os.path.join(PROCESSED_DIR, "val")
test_dir = os.path.join(PROCESSED_DIR, "test")

all_paths = pipeline.get_image_paths(PROCESSED_DIR)
train_paths = [path for path in all_paths if "/train/" in path]
mean, std = pipeline.calc_mean_std(train_paths)
```

```
[4]: print("loading train/val/test generators from data_pipeline")
train_data_gen, val_data_gen, test_data_gen, test_data_gen_raw = pipeline.
    ↪load_data(
        train_dir, val_dir, test_dir, mean, std
```

```
)
```

```
loading train/val/test generators from data_pipeline
creating train generator
Found 1600 images belonging to 2 classes.
creating validation generator
Found 400 images belonging to 2 classes.
creating test generator (normalized)
Found 200 images belonging to 2 classes.
creating test generator (raw)
Found 200 images belonging to 2 classes.
```

```
[5]: # get class names
class_names = list(train_data_gen.class_indices.keys())
print(f"class names found: {class_names}")
```

```
class names found: ['COVID', 'NORMAL']
```

```
[6]: # build and train initial model
input_shape = (IMG_HEIGHT, IMG_WIDTH, NUM_CHANNELS)
model = build_transfer_model(input_shape)
model.summary()
```

```
Model: "functional"
```

Layer (type)	Output Shape	Param #
input_layer_1 ( <a href="#">InputLayer</a> )	( <a href="#">None</a> , 224, 224, 3)	0
resnet50v2 ( <a href="#">Functional</a> )	( <a href="#">None</a> , 7, 7, 2048)	23,564,800
global_average_pooling2d ( <a href="#">GlobalAveragePooling2D</a> )	( <a href="#">None</a> , 2048)	0
dense ( <a href="#">Dense</a> )	( <a href="#">None</a> , 512)	1,049,088
dropout ( <a href="#">Dropout</a> )	( <a href="#">None</a> , 512)	0
dense_1 ( <a href="#">Dense</a> )	( <a href="#">None</a> , 256)	131,328
dropout_1 ( <a href="#">Dropout</a> )	( <a href="#">None</a> , 256)	0
dense_2 ( <a href="#">Dense</a> )	( <a href="#">None</a> , 1)	257

```
Total params: 24,745,473 (94.40 MB)
```

Trainable params: 1,180,673 (4.50 MB)

Non-trainable params: 23,564,800 (89.89 MB)

```
[7]: # train initial model
print("\ntraining initial model")
history = train_model(model, train_data_gen, val_data_gen)
plot_training_history(history, "initial training")
```

training initial model

```
/opt/anaconda3/envs/ml-2025/lib/python3.12/site-
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
```

```
self._warn_if_super_not_called()
```

Epoch 1/30

```
13/13          44s 3s/step -
accuracy: 0.5746 - loss: 1.8392 - precision: 0.5420 - recall: 0.5806 -
val_accuracy: 0.6575 - val_loss: 0.6266 - val_precision: 0.7405 - val_recall:
0.4850
```

Epoch 2/30

```
13/13          40s 3s/step -
accuracy: 0.7120 - loss: 0.6339 - precision: 0.7419 - recall: 0.6422 -
val_accuracy: 0.5425 - val_loss: 0.6876 - val_precision: 0.6667 - val_recall:
0.1700
```

Epoch 3/30

```
13/13          40s 3s/step -
accuracy: 0.7535 - loss: 0.4982 - precision: 0.7486 - recall: 0.7616 -
val_accuracy: 0.6900 - val_loss: 0.6094 - val_precision: 0.7209 - val_recall:
0.6200
```

Epoch 4/30

```
13/13          40s 3s/step -
accuracy: 0.7889 - loss: 0.4372 - precision: 0.7774 - recall: 0.8030 -
val_accuracy: 0.6850 - val_loss: 0.6024 - val_precision: 0.6947 - val_recall:
0.6600
```

Epoch 5/30

```
13/13          40s 3s/step -
accuracy: 0.8163 - loss: 0.3724 - precision: 0.8126 - recall: 0.8260 -
val_accuracy: 0.7125 - val_loss: 0.5839 - val_precision: 0.6809 - val_recall:
0.8000
```

Epoch 6/30

```
13/13          40s 3s/step -
accuracy: 0.8247 - loss: 0.3752 - precision: 0.8081 - recall: 0.8411 -
```

val\_accuracy: 0.7025 - val\_loss: 0.5945 - val\_precision: 0.6431 - val\_recall: 0.9100

Epoch 7/30

13/13 40s 3s/step -

accuracy: 0.8239 - loss: 0.3705 - precision: 0.8146 - recall: 0.8436 -

val\_accuracy: 0.7400 - val\_loss: 0.5486 - val\_precision: 0.7759 - val\_recall: 0.6750

Epoch 8/30

13/13 40s 3s/step -

accuracy: 0.8299 - loss: 0.3608 - precision: 0.8266 - recall: 0.8286 -

val\_accuracy: 0.7675 - val\_loss: 0.5365 - val\_precision: 0.7378 - val\_recall: 0.8300

Epoch 9/30

13/13 40s 3s/step -

accuracy: 0.8346 - loss: 0.3314 - precision: 0.8120 - recall: 0.8741 -

val\_accuracy: 0.7200 - val\_loss: 0.5601 - val\_precision: 0.6606 - val\_recall: 0.9050

Epoch 10/30

13/13 40s 3s/step -

accuracy: 0.8510 - loss: 0.3338 - precision: 0.8252 - recall: 0.8868 -

val\_accuracy: 0.7500 - val\_loss: 0.5220 - val\_precision: 0.7273 - val\_recall: 0.8000

Epoch 11/30

13/13 40s 3s/step -

accuracy: 0.8558 - loss: 0.2989 - precision: 0.8565 - recall: 0.8603 -

val\_accuracy: 0.7775 - val\_loss: 0.5065 - val\_precision: 0.8033 - val\_recall: 0.7350

Epoch 12/30

13/13 40s 3s/step -

accuracy: 0.8499 - loss: 0.3146 - precision: 0.8393 - recall: 0.8658 -

val\_accuracy: 0.7450 - val\_loss: 0.5073 - val\_precision: 0.6788 - val\_recall: 0.9300

Epoch 13/30

13/13 40s 3s/step -

accuracy: 0.8618 - loss: 0.3188 - precision: 0.8311 - recall: 0.9065 -

val\_accuracy: 0.7875 - val\_loss: 0.4807 - val\_precision: 0.7468 - val\_recall: 0.8700

Epoch 14/30

13/13 40s 3s/step -

accuracy: 0.8726 - loss: 0.2979 - precision: 0.8530 - recall: 0.9047 -

val\_accuracy: 0.7650 - val\_loss: 0.4843 - val\_precision: 0.7137 - val\_recall: 0.8850

Epoch 15/30


13/13 40s 3s/step -

accuracy: 0.8695 - loss: 0.2832 - precision: 0.8572 - recall: 0.8805 -

val\_accuracy: 0.7800 - val\_loss: 0.4539 - val\_precision: 0.7240 - val\_recall: 0.9050

Epoch 16/30

13/13                    40s 3s/step -  
accuracy: 0.8630 - loss: 0.3054 - precision: 0.8323 - recall: 0.9100 -  
val\_accuracy: 0.7850 - val\_loss: 0.4508 - val\_precision: 0.7415 - val\_recall:  
0.8750  
Epoch 17/30  
13/13                    40s 3s/step -  
accuracy: 0.8922 - loss: 0.2716 - precision: 0.8699 - recall: 0.9331 -  
val\_accuracy: 0.7900 - val\_loss: 0.4576 - val\_precision: 0.7522 - val\_recall:  
0.8650  
Epoch 18/30  
13/13                    40s 3s/step -  
accuracy: 0.8738 - loss: 0.2698 - precision: 0.8630 - recall: 0.8836 -  
val\_accuracy: 0.7850 - val\_loss: 0.4671 - val\_precision: 0.7395 - val\_recall:  
0.8800  
Epoch 19/30  
13/13                    40s 3s/step -  
accuracy: 0.8857 - loss: 0.2576 - precision: 0.8658 - recall: 0.9031 -  
val\_accuracy: 0.7625 - val\_loss: 0.4985 - val\_precision: 0.6909 - val\_recall:  
0.9500  
Epoch 20/30  
13/13                    40s 3s/step -  
accuracy: 0.8824 - loss: 0.2645 - precision: 0.8571 - recall: 0.9218 -  
val\_accuracy: 0.8150 - val\_loss: 0.4602 - val\_precision: 0.8281 - val\_recall:  
0.7950  
Epoch 21/30  
13/13                    40s 3s/step -  
accuracy: 0.8757 - loss: 0.2740 - precision: 0.8584 - recall: 0.9131 -  
val\_accuracy: 0.7375 - val\_loss: 0.5310 - val\_precision: 0.6727 - val\_recall:  
0.9250



```
[8]: model_save_path = "../models/initial_resnet_model.keras"
      print(f"\nsaving final model to {model_save_path}")
      model.save(model_save_path)
```

saving final model to ../models/initial\_resnet\_model.keras

```
[9]: # hyperparams
      batch_sizes = [16, 32, 64]
      learning_rates = [0.001, 0.0001, 0.00001]
      dropout_rates = [0.2, 0.3, 0.4]

      # param grid
      param_grid = list(itertools.product(batch_sizes, learning_rates, dropout_rates))

      # define number of combinations to randomly sample
      num_combinations_to_test = 5

      # randomly sample combinations
      sampled_params = random.sample(param_grid, num_combinations_to_test)
      print(
          f"randomly sampling {num_combinations_to_test} combinations from_
          ↪{len(param_grid)} total."
      )
```

randomly sampling 5 combinations from 27 total.

```
[10]: # hyperparameter tuning
print("\nperforming hyperparameter tuning")

results = []

for batch_size, learning_rate, dropout_rate in sampled_params:
    print(
        f"training with batch_size={batch_size}, learning_rate={learning_rate},
        ↪ dropout_rate={dropout_rate}"
    )

    # update batch size in generators
    train_data_gen.batch_size = batch_size
    val_data_gen.batch_size = batch_size

    # build and train model
    model = build_transfer_model(input_shape, dropout_rate=dropout_rate)
    history = train_model(
        model, train_data_gen, val_data_gen, learning_rate=learning_rate
    )

    # get best validation metrics
    best_val_loss = min(history.history["val_loss"])
    best_val_acc = max(history.history["val_accuracy"])
    best_val_precision = max(history.history["val_precision"])
    best_val_recall = max(history.history["val_recall"])

    results.append(
        {
            "batch_size": batch_size,
            "learning_rate": learning_rate,
            "dropout_rate": dropout_rate,
            "val_loss": best_val_loss,
            "val_accuracy": best_val_acc,
            "val_precision": best_val_precision,
            "val_recall": best_val_recall,
        }
    )

# save results to csv
os.makedirs("../results", exist_ok=True)
results_df = pd.DataFrame(results)
results_df.to_csv("../results/transfer_learning_hyperparameter_tuning.csv",
    ↪ index=False)
print("\nhyperparameter tuning results:")
```

```
print(results_df)
```

performing hyperparameter tuning

training with batch\_size=64, learning\_rate=1e-05, dropout\_rate=0.2

Epoch 1/30

25/25 46s 2s/step -

accuracy: 0.6027 - loss: 0.9223 - precision: 0.5941 - recall: 0.6153 -  
val\_accuracy: 0.5900 - val\_loss: 0.7056 - val\_precision: 0.5526 - val\_recall:  
0.9450

Epoch 2/30

25/25 41s 2s/step -

accuracy: 0.6149 - loss: 0.7549 - precision: 0.6052 - recall: 0.7115 -  
val\_accuracy: 0.6725 - val\_loss: 0.6292 - val\_precision: 0.6342 - val\_recall:  
0.8150

Epoch 3/30

25/25 42s 2s/step -

accuracy: 0.6776 - loss: 0.6617 - precision: 0.6797 - recall: 0.6827 -  
val\_accuracy: 0.7000 - val\_loss: 0.6072 - val\_precision: 0.6667 - val\_recall:  
0.8000

Epoch 4/30

25/25 41s 2s/step -

accuracy: 0.7293 - loss: 0.5873 - precision: 0.7330 - recall: 0.7266 -  
val\_accuracy: 0.7025 - val\_loss: 0.5994 - val\_precision: 0.6653 - val\_recall:  
0.8150

Epoch 5/30

25/25 41s 2s/step -

accuracy: 0.7440 - loss: 0.5574 - precision: 0.7423 - recall: 0.7283 -  
val\_accuracy: 0.6925 - val\_loss: 0.5942 - val\_precision: 0.6611 - val\_recall:  
0.7900

Epoch 6/30

25/25 41s 2s/step -

accuracy: 0.7308 - loss: 0.5736 - precision: 0.7234 - recall: 0.7541 -  
val\_accuracy: 0.6925 - val\_loss: 0.5946 - val\_precision: 0.6624 - val\_recall:  
0.7850

Epoch 7/30

25/25 41s 2s/step -

accuracy: 0.7159 - loss: 0.5836 - precision: 0.6979 - recall: 0.7342 -  
val\_accuracy: 0.7125 - val\_loss: 0.5866 - val\_precision: 0.6923 - val\_recall:  
0.7650

Epoch 8/30

25/25 41s 2s/step -

accuracy: 0.7366 - loss: 0.5398 - precision: 0.7320 - recall: 0.7271 -  
val\_accuracy: 0.7050 - val\_loss: 0.5862 - val\_precision: 0.6830 - val\_recall:  
0.7650

Epoch 9/30

25/25 41s 2s/step -

accuracy: 0.7587 - loss: 0.4975 - precision: 0.7554 - recall: 0.7327 -



val\_accuracy: 0.6900 - val\_loss: 0.5885 - val\_precision: 0.6583 - val\_recall: 0.7900  
 Epoch 10/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7558 - loss: 0.5159 - precision: 0.7478 - recall: 0.7614 -  
 val\_accuracy: 0.6925 - val\_loss: 0.5820 - val\_precision: 0.6624 - val\_recall: 0.7850  
 Epoch 11/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7685 - loss: 0.4947 - precision: 0.7461 - recall: 0.7842 -  
 val\_accuracy: 0.7075 - val\_loss: 0.5738 - val\_precision: 0.6895 - val\_recall: 0.7550  
 Epoch 12/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7776 - loss: 0.4626 - precision: 0.7653 - recall: 0.7916 -  
 val\_accuracy: 0.7200 - val\_loss: 0.5725 - val\_precision: 0.7056 - val\_recall: 0.7550  
 Epoch 13/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7785 - loss: 0.4639 - precision: 0.7873 - recall: 0.7751 -  
 val\_accuracy: 0.7150 - val\_loss: 0.5748 - val\_precision: 0.6991 - val\_recall: 0.7550  
 Epoch 14/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7770 - loss: 0.4708 - precision: 0.7758 - recall: 0.7693 -  
 val\_accuracy: 0.7000 - val\_loss: 0.5753 - val\_precision: 0.6770 - val\_recall: 0.7650  
 Epoch 15/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7706 - loss: 0.4779 - precision: 0.7612 - recall: 0.7845 -  
 val\_accuracy: 0.7225 - val\_loss: 0.5717 - val\_precision: 0.7129 - val\_recall: 0.7450  
 Epoch 16/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7802 - loss: 0.4386 - precision: 0.7651 - recall: 0.7883 -  
 val\_accuracy: 0.7250 - val\_loss: 0.5663 - val\_precision: 0.7163 - val\_recall: 0.7450  
 Epoch 17/30  
 25/25                    41s 2s/step -  
 accuracy: 0.8053 - loss: 0.4158 - precision: 0.8010 - recall: 0.8115 -  
 val\_accuracy: 0.7275 - val\_loss: 0.5630 - val\_precision: 0.7198 - val\_recall: 0.7450  
 Epoch 18/30  
 25/25                    41s 2s/step -  
 accuracy: 0.8040 - loss: 0.4270 - precision: 0.8071 - recall: 0.8059 -  
 val\_accuracy: 0.7300 - val\_loss: 0.5655 - val\_precision: 0.7130 - val\_recall: 0.7700  
 Epoch 19/30

25/25                    42s 2s/step -  
accuracy: 0.7845 - loss: 0.4395 - precision: 0.7777 - recall: 0.7975 -  
val\_accuracy: 0.7200 - val\_loss: 0.5638 - val\_precision: 0.6982 - val\_recall:  
0.7750  
Epoch 20/30  
25/25                    42s 2s/step -  
accuracy: 0.7929 - loss: 0.4184 - precision: 0.7969 - recall: 0.8103 -  
val\_accuracy: 0.7225 - val\_loss: 0.5586 - val\_precision: 0.7109 - val\_recall:  
0.7500  
Epoch 21/30  
25/25                    42s 2s/step -  
accuracy: 0.7970 - loss: 0.4245 - precision: 0.7858 - recall: 0.7898 -  
val\_accuracy: 0.7200 - val\_loss: 0.5599 - val\_precision: 0.6947 - val\_recall:  
0.7850  
Epoch 22/30  
25/25                    42s 2s/step -  
accuracy: 0.8171 - loss: 0.3994 - precision: 0.7962 - recall: 0.8442 -  
val\_accuracy: 0.7375 - val\_loss: 0.5530 - val\_precision: 0.7209 - val\_recall:  
0.7750  
Epoch 23/30  
25/25                    41s 2s/step -  
accuracy: 0.7813 - loss: 0.4379 - precision: 0.7727 - recall: 0.7843 -  
val\_accuracy: 0.7375 - val\_loss: 0.5509 - val\_precision: 0.7230 - val\_recall:  
0.7700  
Epoch 24/30  
25/25                    41s 2s/step -  
accuracy: 0.8199 - loss: 0.3862 - precision: 0.8197 - recall: 0.8233 -  
val\_accuracy: 0.7400 - val\_loss: 0.5513 - val\_precision: 0.7143 - val\_recall:  
0.8000  
Epoch 25/30  
25/25                    41s 2s/step -  
accuracy: 0.8258 - loss: 0.3926 - precision: 0.8146 - recall: 0.8433 -  
val\_accuracy: 0.7375 - val\_loss: 0.5476 - val\_precision: 0.7251 - val\_recall:  
0.7650  
Epoch 26/30  
25/25                    41s 2s/step -  
accuracy: 0.8109 - loss: 0.4191 - precision: 0.8081 - recall: 0.8089 -  
val\_accuracy: 0.7400 - val\_loss: 0.5475 - val\_precision: 0.7124 - val\_recall:  
0.8050  
Epoch 27/30  
25/25                    42s 2s/step -  
accuracy: 0.8056 - loss: 0.3952 - precision: 0.7946 - recall: 0.8295 -  
val\_accuracy: 0.7450 - val\_loss: 0.5407 - val\_precision: 0.7290 - val\_recall:  
0.7800  
Epoch 28/30  
25/25                    41s 2s/step -  
accuracy: 0.8330 - loss: 0.3761 - precision: 0.8330 - recall: 0.8281 -  
val\_accuracy: 0.7325 - val\_loss: 0.5359 - val\_precision: 0.7123 - val\_recall:

0.7800  
Epoch 29/30  
25/25                   41s 2s/step -  
accuracy: 0.8002 - loss: 0.4161 - precision: 0.7812 - recall: 0.8153 -  
val\_accuracy: 0.7450 - val\_loss: 0.5369 - val\_precision: 0.7112 - val\_recall:  
0.8250  
Epoch 30/30  
25/25                   41s 2s/step -  
accuracy: 0.8294 - loss: 0.3740 - precision: 0.8163 - recall: 0.8568 -  
val\_accuracy: 0.7475 - val\_loss: 0.5311 - val\_precision: 0.7368 - val\_recall:  
0.7700  
training with batch\_size=32, learning\_rate=1e-05, dropout\_rate=0.3  
Epoch 1/30  
50/50                   45s 840ms/step -  
accuracy: 0.6525 - loss: 0.9082 - precision: 0.6616 - recall: 0.6166 -  
val\_accuracy: 0.6450 - val\_loss: 0.7077 - val\_precision: 0.6058 - val\_recall:  
0.8300  
Epoch 2/30  
50/50                   41s 813ms/step -  
accuracy: 0.6721 - loss: 0.7559 - precision: 0.6362 - recall: 0.6919 -  
val\_accuracy: 0.6600 - val\_loss: 0.6804 - val\_precision: 0.6240 - val\_recall:  
0.8050  
Epoch 3/30  
50/50                   41s 828ms/step -  
accuracy: 0.6935 - loss: 0.6813 - precision: 0.6879 - recall: 0.7098 -  
val\_accuracy: 0.6700 - val\_loss: 0.6622 - val\_precision: 0.6417 - val\_recall:  
0.7700  
Epoch 4/30  
50/50                   41s 826ms/step -  
accuracy: 0.6985 - loss: 0.6587 - precision: 0.6940 - recall: 0.7064 -  
val\_accuracy: 0.6575 - val\_loss: 0.6520 - val\_precision: 0.6352 - val\_recall:  
0.7400  
Epoch 5/30  
50/50                   41s 830ms/step -  
accuracy: 0.7439 - loss: 0.5373 - precision: 0.7424 - recall: 0.7510 -  
val\_accuracy: 0.6850 - val\_loss: 0.6395 - val\_precision: 0.6832 - val\_recall:  
0.6900  
Epoch 6/30  
50/50                   41s 820ms/step -  
accuracy: 0.7533 - loss: 0.5558 - precision: 0.7721 - recall: 0.7539 -  
val\_accuracy: 0.6750 - val\_loss: 0.6354 - val\_precision: 0.6699 - val\_recall:  
0.6900  
Epoch 7/30  
50/50                   41s 823ms/step -  
accuracy: 0.7632 - loss: 0.5134 - precision: 0.7598 - recall: 0.7536 -  
val\_accuracy: 0.6750 - val\_loss: 0.6302 - val\_precision: 0.6716 - val\_recall:  
0.6850  
Epoch 8/30

50/50                    41s 822ms/step -  
accuracy: 0.7540 - loss: 0.5294 - precision: 0.7631 - recall: 0.7168 -  
val\_accuracy: 0.6750 - val\_loss: 0.6312 - val\_precision: 0.6549 - val\_recall:  
0.7400  
Epoch 9/30  
50/50                    41s 822ms/step -  
accuracy: 0.7711 - loss: 0.5066 - precision: 0.7722 - recall: 0.7879 -  
val\_accuracy: 0.6800 - val\_loss: 0.6210 - val\_precision: 0.6800 - val\_recall:  
0.6800  
Epoch 10/30  
50/50                    41s 822ms/step -  
accuracy: 0.7697 - loss: 0.5075 - precision: 0.7994 - recall: 0.7464 -  
val\_accuracy: 0.6850 - val\_loss: 0.6132 - val\_precision: 0.6888 - val\_recall:  
0.6750  
Epoch 11/30  
50/50                    41s 819ms/step -  
accuracy: 0.7801 - loss: 0.4798 - precision: 0.7918 - recall: 0.7606 -  
val\_accuracy: 0.6975 - val\_loss: 0.6032 - val\_precision: 0.6854 - val\_recall:  
0.7300  
Epoch 12/30  
50/50                    42s 832ms/step -  
accuracy: 0.7834 - loss: 0.4743 - precision: 0.7648 - recall: 0.7871 -  
val\_accuracy: 0.6875 - val\_loss: 0.5998 - val\_precision: 0.6829 - val\_recall:  
0.7000  
Epoch 13/30  
50/50                    41s 825ms/step -  
accuracy: 0.7939 - loss: 0.4779 - precision: 0.7906 - recall: 0.8031 -  
val\_accuracy: 0.7000 - val\_loss: 0.5957 - val\_precision: 0.6869 - val\_recall:  
0.7350  
Epoch 14/30  
50/50                    41s 829ms/step -  
accuracy: 0.8012 - loss: 0.4600 - precision: 0.8067 - recall: 0.7822 -  
val\_accuracy: 0.7200 - val\_loss: 0.5942 - val\_precision: 0.6982 - val\_recall:  
0.7750  
Epoch 15/30  
50/50                    41s 823ms/step -  
accuracy: 0.8052 - loss: 0.4442 - precision: 0.7989 - recall: 0.8004 -  
val\_accuracy: 0.7050 - val\_loss: 0.5940 - val\_precision: 0.6916 - val\_recall:  
0.7400  
Epoch 16/30  
50/50                    41s 828ms/step -  
accuracy: 0.8124 - loss: 0.4495 - precision: 0.8366 - recall: 0.8122 -  
val\_accuracy: 0.7050 - val\_loss: 0.5929 - val\_precision: 0.7030 - val\_recall:  
0.7100  
Epoch 17/30  
50/50                    41s 819ms/step -  
accuracy: 0.7992 - loss: 0.4258 - precision: 0.8015 - recall: 0.7933 -  
val\_accuracy: 0.7050 - val\_loss: 0.5917 - val\_precision: 0.6952 - val\_recall:

0.7300  
Epoch 18/30  
50/50                   41s 823ms/step -  
accuracy: 0.7931 - loss: 0.4287 - precision: 0.7917 - recall: 0.7893 -  
val\_accuracy: 0.7050 - val\_loss: 0.5855 - val\_precision: 0.6864 - val\_recall:  
0.7550  
Epoch 19/30  
50/50                   41s 819ms/step -  
accuracy: 0.7795 - loss: 0.4543 - precision: 0.7606 - recall: 0.8107 -  
val\_accuracy: 0.7125 - val\_loss: 0.5845 - val\_precision: 0.6977 - val\_recall:  
0.7500  
Epoch 20/30  
50/50                   42s 832ms/step -  
accuracy: 0.8111 - loss: 0.3904 - precision: 0.8029 - recall: 0.8282 -  
val\_accuracy: 0.7150 - val\_loss: 0.5822 - val\_precision: 0.7172 - val\_recall:  
0.7100  
Epoch 21/30  
50/50                   41s 824ms/step -  
accuracy: 0.8154 - loss: 0.4093 - precision: 0.8328 - recall: 0.7891 -  
val\_accuracy: 0.7100 - val\_loss: 0.5800 - val\_precision: 0.7000 - val\_recall:  
0.7350  
Epoch 22/30  
50/50                   41s 830ms/step -  
accuracy: 0.8184 - loss: 0.3951 - precision: 0.8213 - recall: 0.8219 -  
val\_accuracy: 0.7250 - val\_loss: 0.5696 - val\_precision: 0.7009 - val\_recall:  
0.7850  
Epoch 23/30  
50/50                   41s 824ms/step -  
accuracy: 0.8168 - loss: 0.4198 - precision: 0.8063 - recall: 0.8327 -  
val\_accuracy: 0.7225 - val\_loss: 0.5628 - val\_precision: 0.7051 - val\_recall:  
0.7650  
Epoch 24/30  
50/50                   42s 832ms/step -  
accuracy: 0.8079 - loss: 0.4247 - precision: 0.8031 - recall: 0.8023 -  
val\_accuracy: 0.7250 - val\_loss: 0.5630 - val\_precision: 0.6974 - val\_recall:  
0.7950  
Epoch 25/30  
50/50                   41s 830ms/step -  
accuracy: 0.8230 - loss: 0.4007 - precision: 0.8032 - recall: 0.8271 -  
val\_accuracy: 0.7225 - val\_loss: 0.5744 - val\_precision: 0.6894 - val\_recall:  
0.8100  
Epoch 26/30  
50/50                   42s 835ms/step -  
accuracy: 0.8026 - loss: 0.4289 - precision: 0.7979 - recall: 0.8243 -  
val\_accuracy: 0.7175 - val\_loss: 0.5669 - val\_precision: 0.6986 - val\_recall:  
0.7650  
Epoch 27/30  
50/50                   42s 832ms/step -

accuracy: 0.8359 - loss: 0.3709 - precision: 0.8337 - recall: 0.8373 -  
 val\_accuracy: 0.7125 - val\_loss: 0.5643 - val\_precision: 0.6889 - val\_recall:  
 0.7750  
 Epoch 28/30  
 50/50                    41s 827ms/step -  
 accuracy: 0.8374 - loss: 0.3646 - precision: 0.8285 - recall: 0.8574 -  
 val\_accuracy: 0.7125 - val\_loss: 0.5658 - val\_precision: 0.7014 - val\_recall:  
 0.7400  
 training with batch\_size=64, learning\_rate=1e-05, dropout\_rate=0.4  
 Epoch 1/30  
 25/25                    44s 2s/step -  
 accuracy: 0.5705 - loss: 1.2596 - precision: 0.5766 - recall: 0.5516 -  
 val\_accuracy: 0.5425 - val\_loss: 0.7211 - val\_precision: 0.5245 - val\_recall:  
 0.9100  
 Epoch 2/30  
 25/25                    41s 2s/step -  
 accuracy: 0.5398 - loss: 1.0139 - precision: 0.5275 - recall: 0.6244 -  
 val\_accuracy: 0.6125 - val\_loss: 0.7094 - val\_precision: 0.5738 - val\_recall:  
 0.8750  
 Epoch 3/30  
 25/25                    41s 2s/step -  
 accuracy: 0.5900 - loss: 0.9051 - precision: 0.5896 - recall: 0.6268 -  
 val\_accuracy: 0.6725 - val\_loss: 0.6972 - val\_precision: 0.6245 - val\_recall:  
 0.8650  
 Epoch 4/30  
 25/25                    41s 2s/step -  
 accuracy: 0.6265 - loss: 0.8613 - precision: 0.6099 - recall: 0.6406 -  
 val\_accuracy: 0.6950 - val\_loss: 0.6923 - val\_precision: 0.6535 - val\_recall:  
 0.8300  
 Epoch 5/30  
 25/25                    41s 2s/step -  
 accuracy: 0.6650 - loss: 0.8157 - precision: 0.6580 - recall: 0.6804 -  
 val\_accuracy: 0.6925 - val\_loss: 0.6587 - val\_precision: 0.6758 - val\_recall:  
 0.7400  
 Epoch 6/30  
 25/25                    41s 2s/step -  
 accuracy: 0.6754 - loss: 0.7310 - precision: 0.6779 - recall: 0.6645 -  
 val\_accuracy: 0.7000 - val\_loss: 0.6637 - val\_precision: 0.6754 - val\_recall:  
 0.7700  
 Epoch 7/30  
 25/25                    41s 2s/step -  
 accuracy: 0.6722 - loss: 0.7459 - precision: 0.6717 - recall: 0.6640 -  
 val\_accuracy: 0.6900 - val\_loss: 0.6519 - val\_precision: 0.6759 - val\_recall:  
 0.7300  
 Epoch 8/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7194 - loss: 0.6489 - precision: 0.7200 - recall: 0.6928 -  
 val\_accuracy: 0.6975 - val\_loss: 0.6414 - val\_precision: 0.6872 - val\_recall:

0.7250  
Epoch 9/30  
25/25 41s 2s/step -  
accuracy: 0.6853 - loss: 0.6609 - precision: 0.6988 - recall: 0.6629 -  
val\_accuracy: 0.7150 - val\_loss: 0.6398 - val\_precision: 0.6955 - val\_recall:  
0.7650  
Epoch 10/30  
25/25 41s 2s/step -  
accuracy: 0.6903 - loss: 0.7003 - precision: 0.6832 - recall: 0.6665 -  
val\_accuracy: 0.7175 - val\_loss: 0.6362 - val\_precision: 0.6968 - val\_recall:  
0.7700  
Epoch 11/30  
25/25 41s 2s/step -  
accuracy: 0.7056 - loss: 0.6371 - precision: 0.7102 - recall: 0.6994 -  
val\_accuracy: 0.7100 - val\_loss: 0.6208 - val\_precision: 0.7000 - val\_recall:  
0.7350  
Epoch 12/30  
25/25 41s 2s/step -  
accuracy: 0.6967 - loss: 0.6466 - precision: 0.7262 - recall: 0.6546 -  
val\_accuracy: 0.7150 - val\_loss: 0.6220 - val\_precision: 0.6972 - val\_recall:  
0.7600  
Epoch 13/30  
25/25 41s 2s/step -  
accuracy: 0.7202 - loss: 0.6048 - precision: 0.7183 - recall: 0.7223 -  
val\_accuracy: 0.7225 - val\_loss: 0.6142 - val\_precision: 0.7032 - val\_recall:  
0.7700  
Epoch 14/30  
25/25 41s 2s/step -  
accuracy: 0.7166 - loss: 0.6168 - precision: 0.7141 - recall: 0.6747 -  
val\_accuracy: 0.7200 - val\_loss: 0.6223 - val\_precision: 0.6930 - val\_recall:  
0.7900  
Epoch 15/30  
25/25 41s 2s/step -  
accuracy: 0.7337 - loss: 0.6058 - precision: 0.7457 - recall: 0.7016 -  
val\_accuracy: 0.7250 - val\_loss: 0.6331 - val\_precision: 0.6875 - val\_recall:  
0.8250  
Epoch 16/30  
25/25 41s 2s/step -  
accuracy: 0.7427 - loss: 0.5732 - precision: 0.7324 - recall: 0.7690 -  
val\_accuracy: 0.7200 - val\_loss: 0.6170 - val\_precision: 0.6930 - val\_recall:  
0.7900  
Epoch 17/30  
25/25 41s 2s/step -  
accuracy: 0.7445 - loss: 0.5458 - precision: 0.7661 - recall: 0.7174 -  
val\_accuracy: 0.7275 - val\_loss: 0.6088 - val\_precision: 0.7059 - val\_recall:  
0.7800  
Epoch 18/30  
25/25 41s 2s/step -

accuracy: 0.7448 - loss: 0.5565 - precision: 0.7599 - recall: 0.7101 -  
 val\_accuracy: 0.7225 - val\_loss: 0.5945 - val\_precision: 0.7032 - val\_recall:  
 0.7700  
 Epoch 19/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7485 - loss: 0.5710 - precision: 0.7640 - recall: 0.7302 -  
 val\_accuracy: 0.7175 - val\_loss: 0.5902 - val\_precision: 0.6968 - val\_recall:  
 0.7700  
 Epoch 20/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7505 - loss: 0.5184 - precision: 0.7460 - recall: 0.7218 -  
 val\_accuracy: 0.7225 - val\_loss: 0.5945 - val\_precision: 0.6926 - val\_recall:  
 0.8000  
 Epoch 21/30  
 25/25                    42s 2s/step -  
 accuracy: 0.7648 - loss: 0.5251 - precision: 0.7664 - recall: 0.7549 -  
 val\_accuracy: 0.7175 - val\_loss: 0.5948 - val\_precision: 0.6883 - val\_recall:  
 0.7950  
 Epoch 22/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7623 - loss: 0.5294 - precision: 0.7819 - recall: 0.7599 -  
 val\_accuracy: 0.7225 - val\_loss: 0.5897 - val\_precision: 0.6996 - val\_recall:  
 0.7800  
 Epoch 23/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7704 - loss: 0.4967 - precision: 0.7869 - recall: 0.7535 -  
 val\_accuracy: 0.7125 - val\_loss: 0.5842 - val\_precision: 0.7033 - val\_recall:  
 0.7350  
 Epoch 24/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7432 - loss: 0.5461 - precision: 0.7555 - recall: 0.7239 -  
 val\_accuracy: 0.7125 - val\_loss: 0.5815 - val\_precision: 0.6941 - val\_recall:  
 0.7600  
 Epoch 25/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7647 - loss: 0.5223 - precision: 0.7655 - recall: 0.7528 -  
 val\_accuracy: 0.7125 - val\_loss: 0.5808 - val\_precision: 0.6959 - val\_recall:  
 0.7550  
 Epoch 26/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7660 - loss: 0.4932 - precision: 0.7636 - recall: 0.7598 -  
 val\_accuracy: 0.7150 - val\_loss: 0.5792 - val\_precision: 0.6972 - val\_recall:  
 0.7600  
 Epoch 27/30  
 25/25                    41s 2s/step -  
 accuracy: 0.7721 - loss: 0.4901 - precision: 0.7749 - recall: 0.7645 -  
 val\_accuracy: 0.7100 - val\_loss: 0.5761 - val\_precision: 0.6892 - val\_recall:  
 0.7650



Epoch 28/30  
 25/25 41s 2s/step -  
 accuracy: 0.7815 - loss: 0.4752 - precision: 0.7784 - recall: 0.7633 -  
 val\_accuracy: 0.7200 - val\_loss: 0.5704 - val\_precision: 0.6947 - val\_recall:  
 0.7850

Epoch 29/30  
 25/25 41s 2s/step -  
 accuracy: 0.7597 - loss: 0.4708 - precision: 0.7539 - recall: 0.7488 -  
 val\_accuracy: 0.7225 - val\_loss: 0.5692 - val\_precision: 0.6960 - val\_recall:  
 0.7900

Epoch 30/30  
 25/25 41s 2s/step -  
 accuracy: 0.7935 - loss: 0.4663 - precision: 0.7949 - recall: 0.7875 -  
 val\_accuracy: 0.7175 - val\_loss: 0.5680 - val\_precision: 0.6916 - val\_recall:  
 0.7850  
 training with batch\_size=32, learning\_rate=0.0001, dropout\_rate=0.3

Epoch 1/30  
 50/50 45s 842ms/step -  
 accuracy: 0.6883 - loss: 0.7471 - precision: 0.6648 - recall: 0.7104 -  
 val\_accuracy: 0.6950 - val\_loss: 0.6174 - val\_precision: 0.6711 - val\_recall:  
 0.7650

Epoch 2/30  
 50/50 41s 830ms/step -  
 accuracy: 0.7583 - loss: 0.5230 - precision: 0.7423 - recall: 0.7739 -  
 val\_accuracy: 0.6975 - val\_loss: 0.6502 - val\_precision: 0.6348 - val\_recall:  
 0.9300

Epoch 3/30  
 50/50 41s 825ms/step -  
 accuracy: 0.7810 - loss: 0.4645 - precision: 0.7656 - recall: 0.8311 -  
 val\_accuracy: 0.7475 - val\_loss: 0.5433 - val\_precision: 0.7037 - val\_recall:  
 0.8550

Epoch 4/30  
 50/50 41s 830ms/step -  
 accuracy: 0.8210 - loss: 0.4157 - precision: 0.8168 - recall: 0.8538 -  
 val\_accuracy: 0.7600 - val\_loss: 0.5595 - val\_precision: 0.7261 - val\_recall:  
 0.8350

Epoch 5/30  
 50/50 41s 825ms/step -  
 accuracy: 0.8065 - loss: 0.4498 - precision: 0.8073 - recall: 0.8225 -  
 val\_accuracy: 0.6575 - val\_loss: 0.6029 - val\_precision: 0.7838 - val\_recall:  
 0.4350

Epoch 6/30  
 50/50 41s 830ms/step -  
 accuracy: 0.8145 - loss: 0.3920 - precision: 0.8096 - recall: 0.8076 -  
 val\_accuracy: 0.7200 - val\_loss: 0.6051 - val\_precision: 0.6571 - val\_recall:  
 0.9200

Epoch 7/30  
 50/50 41s 828ms/step -

accuracy: 0.8453 - loss: 0.3553 - precision: 0.8253 - recall: 0.8797 -  
 val\_accuracy: 0.7550 - val\_loss: 0.4994 - val\_precision: 0.7125 - val\_recall: 0.8550  
 Epoch 8/30  
 50/50 41s 821ms/step -  
 accuracy: 0.8547 - loss: 0.3413 - precision: 0.8453 - recall: 0.8708 -  
 val\_accuracy: 0.7250 - val\_loss: 0.5477 - val\_precision: 0.6718 - val\_recall: 0.8800  
 Epoch 9/30  
 50/50 41s 819ms/step -  
 accuracy: 0.8556 - loss: 0.3183 - precision: 0.8473 - recall: 0.8651 -  
 val\_accuracy: 0.7750 - val\_loss: 0.4795 - val\_precision: 0.7292 - val\_recall: 0.8750  
 Epoch 10/30  
 50/50 41s 824ms/step -  
 accuracy: 0.8536 - loss: 0.3520 - precision: 0.8458 - recall: 0.8669 -  
 val\_accuracy: 0.7825 - val\_loss: 0.4870 - val\_precision: 0.7534 - val\_recall: 0.8400  
 Epoch 11/30  
 50/50 41s 825ms/step -  
 accuracy: 0.8463 - loss: 0.3266 - precision: 0.8297 - recall: 0.8598 -  
 val\_accuracy: 0.7800 - val\_loss: 0.4745 - val\_precision: 0.7917 - val\_recall: 0.7600  
 Epoch 12/30  
 50/50 41s 826ms/step -  
 accuracy: 0.8781 - loss: 0.2925 - precision: 0.8654 - recall: 0.9057 -  
 val\_accuracy: 0.7550 - val\_loss: 0.4999 - val\_precision: 0.6932 - val\_recall: 0.9150  
 Epoch 13/30  
 50/50 41s 829ms/step -  
 accuracy: 0.8593 - loss: 0.3107 - precision: 0.8363 - recall: 0.8881 -  
 val\_accuracy: 0.7800 - val\_loss: 0.4649 - val\_precision: 0.7545 - val\_recall: 0.8300  
 Epoch 14/30  
 50/50 41s 825ms/step -  
 accuracy: 0.8650 - loss: 0.2893 - precision: 0.8544 - recall: 0.8824 -  
 val\_accuracy: 0.7300 - val\_loss: 0.5311 - val\_precision: 0.6631 - val\_recall: 0.9350  
 Epoch 15/30  
 50/50 41s 820ms/step -  
 accuracy: 0.8714 - loss: 0.2932 - precision: 0.8532 - recall: 0.9016 -  
 val\_accuracy: 0.7850 - val\_loss: 0.4838 - val\_precision: 0.7938 - val\_recall: 0.7700  
 Epoch 16/30  
 50/50 41s 822ms/step -  
 accuracy: 0.8744 - loss: 0.2796 - precision: 0.8696 - recall: 0.8838 -  
 val\_accuracy: 0.7925 - val\_loss: 0.4584 - val\_precision: 0.7555 - val\_recall: 0.8650

Epoch 17/30  
50/50 41s 821ms/step -  
accuracy: 0.8708 - loss: 0.2845 - precision: 0.8623 - recall: 0.8941 -  
val\_accuracy: 0.7925 - val\_loss: 0.4463 - val\_precision: 0.7532 - val\_recall:  
0.8700  
Epoch 18/30  
50/50 41s 823ms/step -  
accuracy: 0.8726 - loss: 0.2847 - precision: 0.8653 - recall: 0.8812 -  
val\_accuracy: 0.7225 - val\_loss: 0.5818 - val\_precision: 0.6540 - val\_recall:  
0.9450  
Epoch 19/30  
50/50 41s 823ms/step -  
accuracy: 0.8698 - loss: 0.2949 - precision: 0.8553 - recall: 0.8993 -  
val\_accuracy: 0.7850 - val\_loss: 0.4444 - val\_precision: 0.7436 - val\_recall:  
0.8700  
Epoch 20/30  
50/50 41s 824ms/step -  
accuracy: 0.8863 - loss: 0.2558 - precision: 0.8749 - recall: 0.9000 -  
val\_accuracy: 0.7450 - val\_loss: 0.4935 - val\_precision: 0.6750 - val\_recall:  
0.9450  
Epoch 21/30  
50/50 41s 824ms/step -  
accuracy: 0.8846 - loss: 0.2738 - precision: 0.8579 - recall: 0.9196 -  
val\_accuracy: 0.7475 - val\_loss: 0.4817 - val\_precision: 0.6800 - val\_recall:  
0.9350  
Epoch 22/30  
50/50 41s 824ms/step -  
accuracy: 0.8865 - loss: 0.2618 - precision: 0.8775 - recall: 0.9020 -  
val\_accuracy: 0.8050 - val\_loss: 0.4368 - val\_precision: 0.7961 - val\_recall:  
0.8200  
Epoch 23/30  
50/50 41s 825ms/step -  
accuracy: 0.9078 - loss: 0.2359 - precision: 0.8917 - recall: 0.9295 -  
val\_accuracy: 0.8025 - val\_loss: 0.4490 - val\_precision: 0.7738 - val\_recall:  
0.8550  
Epoch 24/30  
50/50 41s 823ms/step -  
accuracy: 0.8927 - loss: 0.2430 - precision: 0.8882 - recall: 0.9010 -  
val\_accuracy: 0.7850 - val\_loss: 0.4825 - val\_precision: 0.8608 - val\_recall:  
0.6800  
Epoch 25/30  
50/50 41s 826ms/step -  
accuracy: 0.8878 - loss: 0.2739 - precision: 0.9017 - recall: 0.8785 -  
val\_accuracy: 0.7725 - val\_loss: 0.4887 - val\_precision: 0.7121 - val\_recall:  
0.9150  
Epoch 26/30  
50/50 41s 822ms/step -  
accuracy: 0.8823 - loss: 0.2572 - precision: 0.8464 - recall: 0.9223 -

val\_accuracy: 0.7525 - val\_loss: 0.4766 - val\_precision: 0.6996 - val\_recall: 0.8850

Epoch 27/30

50/50 41s 828ms/step -

accuracy: 0.8965 - loss: 0.2372 - precision: 0.8802 - recall: 0.9224 -

val\_accuracy: 0.8375 - val\_loss: 0.4194 - val\_precision: 0.8462 - val\_recall: 0.8250

Epoch 28/30

50/50 41s 825ms/step -

accuracy: 0.8957 - loss: 0.2254 - precision: 0.8797 - recall: 0.9242 -

val\_accuracy: 0.7575 - val\_loss: 0.4930 - val\_precision: 0.6900 - val\_recall: 0.9350

Epoch 29/30

50/50 41s 825ms/step -

accuracy: 0.8842 - loss: 0.2540 - precision: 0.8785 - recall: 0.8969 -

val\_accuracy: 0.8300 - val\_loss: 0.3942 - val\_precision: 0.8333 - val\_recall: 0.8250

Epoch 30/30

50/50 41s 827ms/step -

accuracy: 0.8919 - loss: 0.2357 - precision: 0.8927 - recall: 0.8913 -

val\_accuracy: 0.7975 - val\_loss: 0.4320 - val\_precision: 0.7960 - val\_recall: 0.8000

training with batch\_size=64, learning\_rate=0.001, dropout\_rate=0.2

Epoch 1/30

25/25 44s 2s/step -

accuracy: 0.7352 - loss: 0.8453 - precision: 0.7387 - recall: 0.7332 -

val\_accuracy: 0.6825 - val\_loss: 0.6150 - val\_precision: 0.6347 - val\_recall: 0.8600

Epoch 2/30

25/25 41s 2s/step -

accuracy: 0.8088 - loss: 0.4031 - precision: 0.7988 - recall: 0.8347 -

val\_accuracy: 0.6975 - val\_loss: 0.5885 - val\_precision: 0.6561 - val\_recall: 0.8300

Epoch 3/30

25/25 41s 2s/step -

accuracy: 0.8098 - loss: 0.3745 - precision: 0.8146 - recall: 0.7957 -

val\_accuracy: 0.7375 - val\_loss: 0.5575 - val\_precision: 0.6834 - val\_recall: 0.8850

Epoch 4/30

25/25 41s 2s/step -

accuracy: 0.8657 - loss: 0.3262 - precision: 0.8454 - recall: 0.9080 -

val\_accuracy: 0.7275 - val\_loss: 0.5572 - val\_precision: 0.6717 - val\_recall: 0.8900

Epoch 5/30

25/25 41s 2s/step -

accuracy: 0.8484 - loss: 0.3398 - precision: 0.8406 - recall: 0.8773 -

val\_accuracy: 0.7400 - val\_loss: 0.5347 - val\_precision: 0.7222 - val\_recall: 0.7800

Epoch 6/30  
25/25 41s 2s/step -  
accuracy: 0.8585 - loss: 0.3141 - precision: 0.8532 - recall: 0.8600 -  
val\_accuracy: 0.7425 - val\_loss: 0.5456 - val\_precision: 0.6830 - val\_recall:  
0.9050  
Epoch 7/30  
25/25 41s 2s/step -  
accuracy: 0.8560 - loss: 0.3354 - precision: 0.8527 - recall: 0.8655 -  
val\_accuracy: 0.6575 - val\_loss: 0.6152 - val\_precision: 0.7692 - val\_recall:  
0.4500  
Epoch 8/30  
25/25 41s 2s/step -  
accuracy: 0.8517 - loss: 0.3017 - precision: 0.8473 - recall: 0.8659 -  
val\_accuracy: 0.7950 - val\_loss: 0.4556 - val\_precision: 0.7837 - val\_recall:  
0.8150  
Epoch 9/30  
25/25 41s 2s/step -  
accuracy: 0.8615 - loss: 0.3043 - precision: 0.8500 - recall: 0.8882 -  
val\_accuracy: 0.8000 - val\_loss: 0.4576 - val\_precision: 0.8191 - val\_recall:  
0.7700  
Epoch 10/30  
25/25 41s 2s/step -  
accuracy: 0.8558 - loss: 0.2978 - precision: 0.8483 - recall: 0.8722 -  
val\_accuracy: 0.7950 - val\_loss: 0.4878 - val\_precision: 0.8687 - val\_recall:  
0.6950  
Epoch 11/30  
25/25 41s 2s/step -  
accuracy: 0.8602 - loss: 0.3234 - precision: 0.8574 - recall: 0.8738 -  
val\_accuracy: 0.7475 - val\_loss: 0.5199 - val\_precision: 0.7438 - val\_recall:  
0.7550  
Epoch 12/30  
25/25 41s 2s/step -  
accuracy: 0.8751 - loss: 0.2766 - precision: 0.8763 - recall: 0.8747 -  
val\_accuracy: 0.7825 - val\_loss: 0.4548 - val\_precision: 0.7181 - val\_recall:  
0.9300  
Epoch 13/30  
25/25 41s 2s/step -  
accuracy: 0.8750 - loss: 0.2612 - precision: 0.8522 - recall: 0.9018 -  
val\_accuracy: 0.8050 - val\_loss: 0.4270 - val\_precision: 0.7990 - val\_recall:  
0.8150  
Epoch 14/30  
25/25 41s 2s/step -  
accuracy: 0.9081 - loss: 0.2353 - precision: 0.9077 - recall: 0.9087 -  
val\_accuracy: 0.7950 - val\_loss: 0.4645 - val\_precision: 0.8278 - val\_recall:  
0.7450  
Epoch 15/30  
25/25 41s 2s/step -  
accuracy: 0.8823 - loss: 0.2713 - precision: 0.8812 - recall: 0.8832 -

val\_accuracy: 0.7625 - val\_loss: 0.5278 - val\_precision: 0.6882 - val\_recall: 0.9600

Epoch 16/30

25/25 41s 2s/step -

accuracy: 0.8899 - loss: 0.2703 - precision: 0.8711 - recall: 0.9170 -

val\_accuracy: 0.7950 - val\_loss: 0.4646 - val\_precision: 0.8315 - val\_recall: 0.7400

Epoch 17/30

25/25 41s 2s/step -

accuracy: 0.8774 - loss: 0.2797 - precision: 0.8774 - recall: 0.8730 -

val\_accuracy: 0.7675 - val\_loss: 0.5150 - val\_precision: 0.7019 - val\_recall: 0.9300

Epoch 18/30

25/25 41s 2s/step -

accuracy: 0.8838 - loss: 0.2471 - precision: 0.8705 - recall: 0.8924 -

val\_accuracy: 0.7575 - val\_loss: 0.4601 - val\_precision: 0.6886 - val\_recall: 0.9400

hyperparameter tuning results:

	batch_size	learning_rate	dropout_rate	val_loss	val_accuracy \
0	64	0.00001	0.2	0.531070	0.7475
1	32	0.00001	0.3	0.562782	0.7250
2	64	0.00001	0.4	0.568049	0.7275
3	32	0.00010	0.3	0.394171	0.8375
4	64	0.00100	0.2	0.426978	0.8050

	val_precision	val_recall
0	0.736842	0.945
1	0.717172	0.830
2	0.705882	0.910
3	0.860759	0.945
4	0.868750	0.960

```
[11]: # concatenate train and val data
X_all = []
y_all = []

train_data_gen.reset()
val_data_gen.reset()

for batch_x, batch_y in train_data_gen:
    X_all.append(batch_x)
    y_all.append(batch_y)
    if len(X_all) * BATCH_SIZE >= train_data_gen.samples:
        break

for batch_x, batch_y in val_data_gen:
```

```

X_all.append(batch_x)
y_all.append(batch_y)
if len(X_all) * BATCH_SIZE >= val_data_gen.samples + train_data_gen.samples:
    break

X_all = np.concatenate(X_all)
y_all = np.concatenate(y_all)

# create dataset
train_val_ds = (
    tf.data.Dataset.from_tensor_slices((X_all, y_all))
    .shuffle(1000)
    .batch(BATCH_SIZE)
    .prefetch(tf.data.AUTOTUNE)
)

```

```

[12]: def plot_training_history(history, title):
    """
    plot training and validation metrics

    params
    -----
    history: tf.keras.callbacks.History
        training history
    title: str
        plot title
    """

    metrics = ["loss", "accuracy", "precision", "recall"]
    fig, axes = plt.subplots(2, 2, figsize=(15, 10))
    axes = axes.ravel()

    for idx, metric in enumerate(metrics):
        axes[idx].plot(history.history[metric], label="train")
        if f"val_{metric}" in history.history:
            axes[idx].plot(history.history[f"val_{metric}"], label="val")
        axes[idx].set_title(f"{metric}")
        axes[idx].set_xlabel("epoch")
        axes[idx].set_ylabel(metric)
        axes[idx].legend()

    plt.suptitle(title)
    plt.tight_layout()
    plt.show()

```

```

[13]: # train final model with best hyperparameters
best_params = results_df.loc[results_df["val_loss"].idxmin()]
print(f"\nbest hyperparameters: {best_params}")

```

```

# update batch size
train_data_gen.batch_size = int(best_params["batch_size"])
val_data_gen.batch_size = int(best_params["batch_size"])

# build and train model
print("\ntraining final model")
model = build_transfer_model(input_shape,
    dropout_rate=best_params["dropout_rate"])
history = train_model(
    model, train_val_ds, None, learning_rate=best_params["learning_rate"]
)
plot_training_history(history, "final training")

# save final model
model_save_path = "../models/final_resnet_model.keras"
print(f"\nsaving final model to {model_save_path}")
model.save(model_save_path)

```

```

best hyperparameters: batch_size      32.000000
learning_rate      0.000100
dropout_rate      0.300000
val_loss          0.394171
val_accuracy      0.837500
val_precision      0.860759
val_recall        0.945000
Name: 3, dtype: float64

```

training final model

Epoch 1/30

125/125 82s 630ms/step -

accuracy: 0.7215 - loss: 0.6307 - precision: 0.7051 - recall: 0.7656

Epoch 2/30

/opt/anaconda3/envs/ml-2025/lib/python3.12/site-

packages/keras/src/callbacks/early\_stopping.py:153: UserWarning: Early stopping  
conditioned on metric `val\_loss` which is not available. Available metrics are:  
accuracy,loss,precision,recall

current = self.get\_monitor\_value(logs)

125/125 78s 627ms/step -

accuracy: 0.7969 - loss: 0.4319 - precision: 0.7876 - recall: 0.8019

Epoch 3/30

125/125 79s 631ms/step -

accuracy: 0.8249 - loss: 0.3747 - precision: 0.8072 - recall: 0.8516

Epoch 4/30

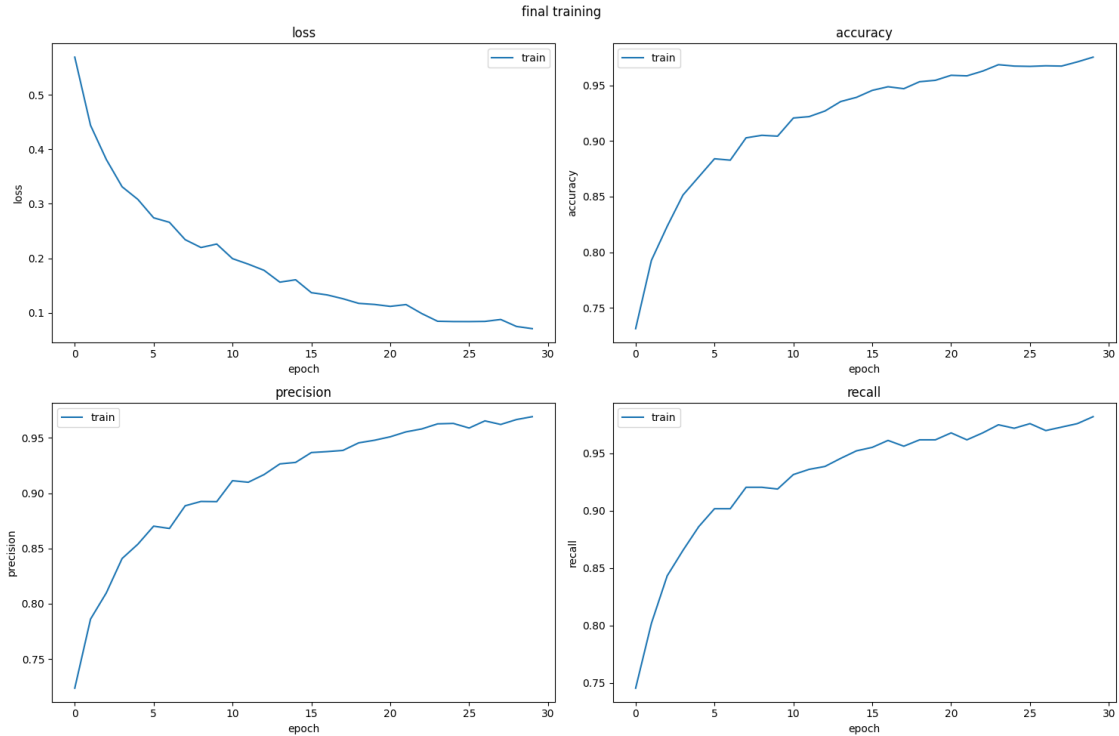
125/125 80s 637ms/step -

accuracy: 0.8620 - loss: 0.3078 - precision: 0.8543 - recall: 0.8737



Epoch 5/30  
125/125 80s 638ms/step -  
accuracy: 0.8763 - loss: 0.2925 - precision: 0.8595 - recall: 0.9029  
Epoch 6/30  
125/125 79s 636ms/step -  
accuracy: 0.8843 - loss: 0.2646 - precision: 0.8702 - recall: 0.9010  
Epoch 7/30  
125/125 79s 635ms/step -  
accuracy: 0.8959 - loss: 0.2510 - precision: 0.8892 - recall: 0.9023  
Epoch 8/30  
125/125 79s 634ms/step -  
accuracy: 0.9099 - loss: 0.2229 - precision: 0.8977 - recall: 0.9269  
Epoch 9/30  
125/125 79s 636ms/step -  
accuracy: 0.9079 - loss: 0.2067 - precision: 0.8885 - recall: 0.9295  
Epoch 10/30  
125/125 81s 646ms/step -  
accuracy: 0.9084 - loss: 0.2085 - precision: 0.8953 - recall: 0.9211  
Epoch 11/30  
125/125 79s 630ms/step -  
accuracy: 0.9291 - loss: 0.1796 - precision: 0.9220 - recall: 0.9359  
Epoch 12/30  
125/125 79s 631ms/step -  
accuracy: 0.9258 - loss: 0.1816 - precision: 0.9138 - recall: 0.9416  
Epoch 13/30  
125/125 79s 631ms/step -  
accuracy: 0.9285 - loss: 0.1679 - precision: 0.9245 - recall: 0.9326  
Epoch 14/30  
125/125 79s 631ms/step -  
accuracy: 0.9339 - loss: 0.1580 - precision: 0.9229 - recall: 0.9447  
Epoch 15/30  
125/125 79s 634ms/step -  
accuracy: 0.9493 - loss: 0.1473 - precision: 0.9355 - recall: 0.9641  
Epoch 16/30  
125/125 79s 629ms/step -  
accuracy: 0.9473 - loss: 0.1335 - precision: 0.9403 - recall: 0.9535  
Epoch 17/30  
125/125 79s 631ms/step -  
accuracy: 0.9517 - loss: 0.1296 - precision: 0.9375 - recall: 0.9661  
Epoch 18/30  
125/125 79s 629ms/step -  
accuracy: 0.9426 - loss: 0.1281 - precision: 0.9374 - recall: 0.9471  
Epoch 19/30  
125/125 79s 630ms/step -  
accuracy: 0.9535 - loss: 0.1155 - precision: 0.9430 - recall: 0.9637  
Epoch 20/30  
125/125 79s 630ms/step -  
accuracy: 0.9555 - loss: 0.1140 - precision: 0.9449 - recall: 0.9661

Epoch 21/30  
125/125 79s 631ms/step -  
accuracy: 0.9615 - loss: 0.1062 - precision: 0.9530 - recall: 0.9707  
Epoch 22/30  
125/125 79s 634ms/step -  
accuracy: 0.9546 - loss: 0.1214 - precision: 0.9533 - recall: 0.9545  
Epoch 23/30  
125/125 79s 631ms/step -  
accuracy: 0.9608 - loss: 0.0963 - precision: 0.9577 - recall: 0.9620  
Epoch 24/30  
125/125 79s 630ms/step -  
accuracy: 0.9714 - loss: 0.0852 - precision: 0.9624 - recall: 0.9814  
Epoch 25/30  
125/125 79s 631ms/step -  
accuracy: 0.9692 - loss: 0.0797 - precision: 0.9631 - recall: 0.9760  
Epoch 26/30  
125/125 79s 631ms/step -  
accuracy: 0.9688 - loss: 0.0809 - precision: 0.9590 - recall: 0.9792  
Epoch 27/30  
125/125 79s 630ms/step -  
accuracy: 0.9653 - loss: 0.0799 - precision: 0.9647 - recall: 0.9645  
Epoch 28/30  
125/125 79s 631ms/step -  
accuracy: 0.9667 - loss: 0.0871 - precision: 0.9571 - recall: 0.9760  
Epoch 29/30  
125/125 79s 631ms/step -  
accuracy: 0.9717 - loss: 0.0720 - precision: 0.9660 - recall: 0.9778  
Epoch 30/30  
125/125 79s 630ms/step -  
accuracy: 0.9773 - loss: 0.0655 - precision: 0.9713 - recall: 0.9830



saving final model to ../models/final\_resnet\_model.keras

```
[14]: # fine-tune model
print("\nfine-tuning")
history_fine = fine_tune_model(model, train_data_gen, val_data_gen)
plot_training_history(history_fine, "fine-tuning")

# save fine-tuned model
model_save_path = "../models/fine_tuned_resnet_model.keras"
print(f"\nsaving fine-tuned model to {model_save_path}")
model.save(model_save_path)
```

fine-tuning

Epoch 1/30

50/50 44s 835ms/step -

accuracy: 0.9661 - loss: 0.2402 - precision: 0.9597 - recall: 0.9730 -

val\_accuracy: 1.0000 - val\_loss: 0.0180 - val\_precision: 1.0000 - val\_recall:

1.0000

Epoch 2/30

50/50 41s 814ms/step -

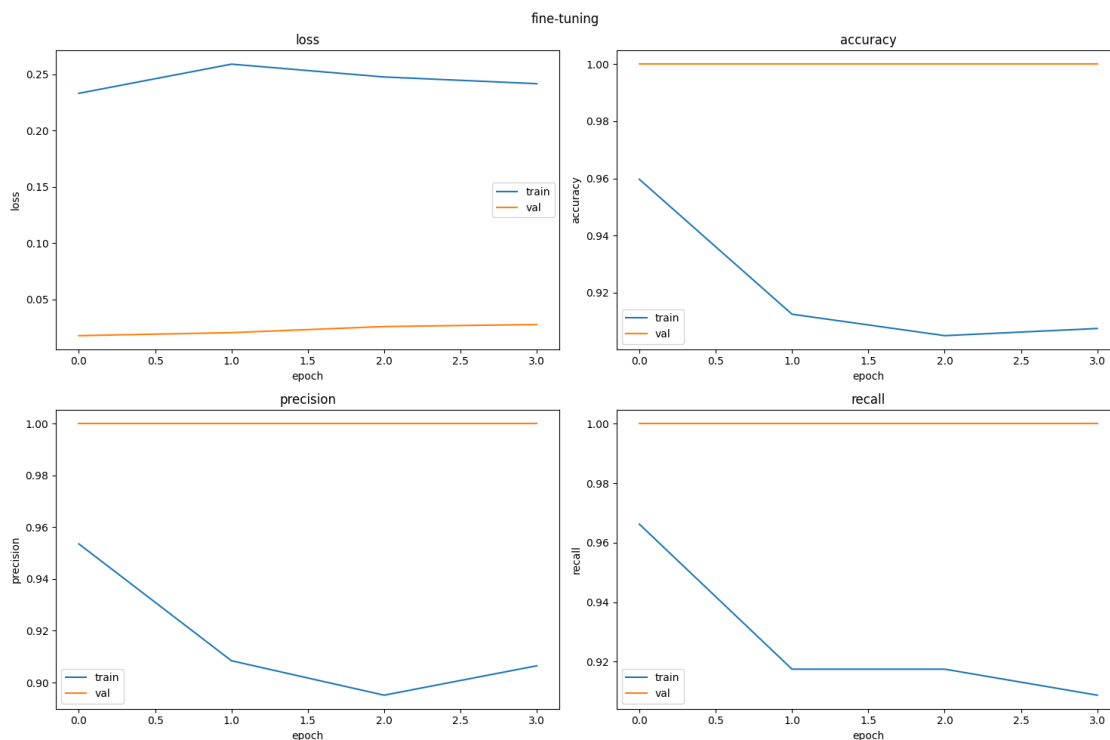
accuracy: 0.9169 - loss: 0.2435 - precision: 0.9043 - recall: 0.9250 -

val\_accuracy: 1.0000 - val\_loss: 0.0206 - val\_precision: 1.0000 - val\_recall:

```

1.0000
Epoch 3/30
50/50          41s 815ms/step -
accuracy: 0.9153 - loss: 0.2247 - precision: 0.9120 - recall: 0.9224 -
val_accuracy: 1.0000 - val_loss: 0.0260 - val_precision: 1.0000 - val_recall:
1.0000
Epoch 4/30
50/50          41s 823ms/step -
accuracy: 0.9041 - loss: 0.2531 - precision: 0.9057 - recall: 0.9015 -
val_accuracy: 1.0000 - val_loss: 0.0278 - val_precision: 1.0000 - val_recall:
1.0000

```



saving fine-tuned model to ../models/fine\_tuned\_resnet\_model.keras

```

[15]: # evaluate on test set
print("\nevaluating on test set")
test_loss, test_acc, test_precision, test_recall = model.evaluate(test_data_gen)
print(f"test loss: {test_loss:.4f}")
print(f"test accuracy: {test_acc:.4f}")
print(f"test precision: {test_precision:.4f}")
print(f"test recall: {test_recall:.4f}")

```

evaluating on test set

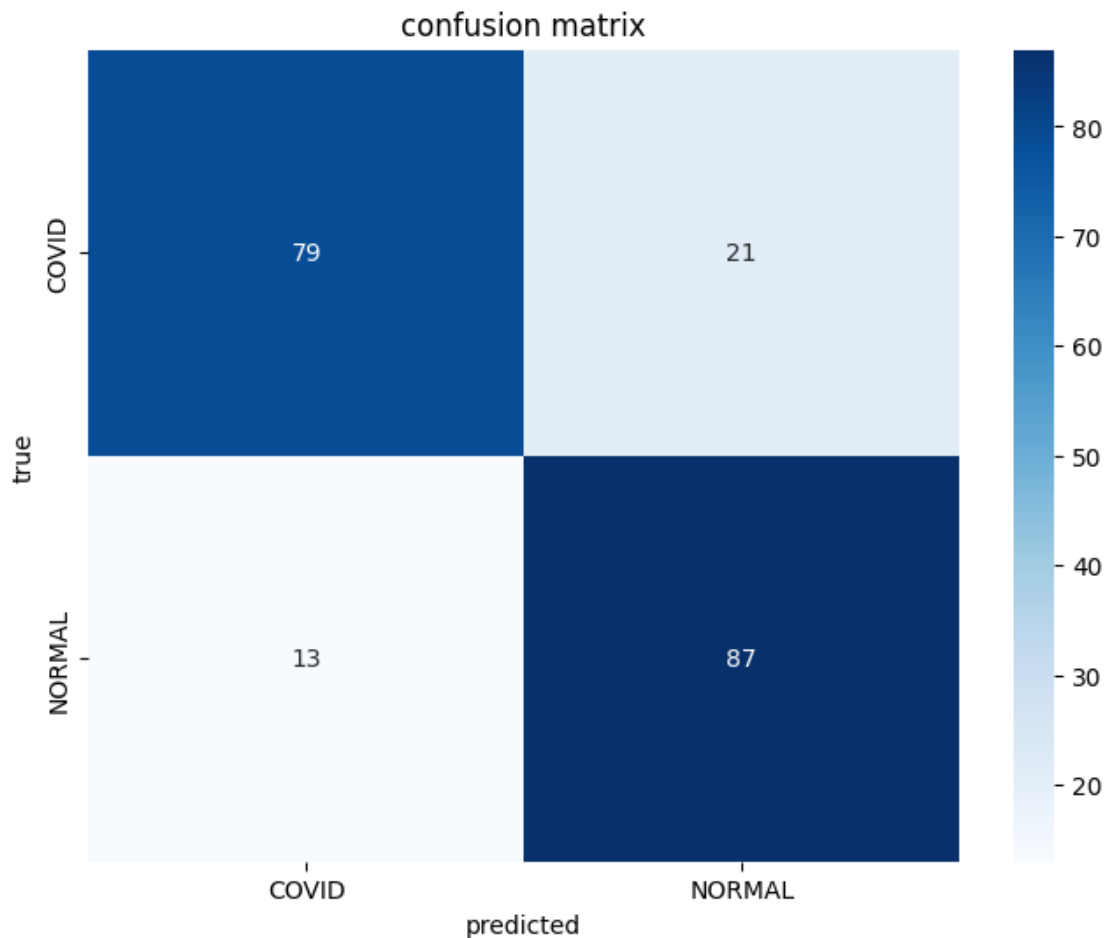
```
/opt/anaconda3/envs/ml-2025/lib/python3.12/site-  
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:  
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in  
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,  
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be  
ignored.
```

```
self._warn_if_super_not_called()
```

```
2/2          4s 2s/step -  
accuracy: 0.8190 - loss: 0.5796 - precision: 0.7113 - recall: 0.8538  
test loss: 0.5247  
test accuracy: 0.8300  
test precision: 0.8056  
test recall: 0.8700
```

```
[16]: # plot confusion matrix  
print("\ngenerating confusion matrix")  
y_pred = model.predict(test_data_gen)  
y_pred = (y_pred > 0.5).astype(int)  
y_true = test_data_gen.classes  
  
plot_confusion_matrix(y_true, y_pred, class_names)
```

```
generating confusion matrix  
2/2          6s 2s/step
```



```
[17]: print("\nplotting sample predictions with raw images")

# get a batch of raw (unnormalized) images and labels
images_raw, labels_raw = next(iter(test_data_gen_raw))

# get the corresponding normalized batch for prediction
test_data_gen.reset()
images_norm, _ = next(iter(test_data_gen))

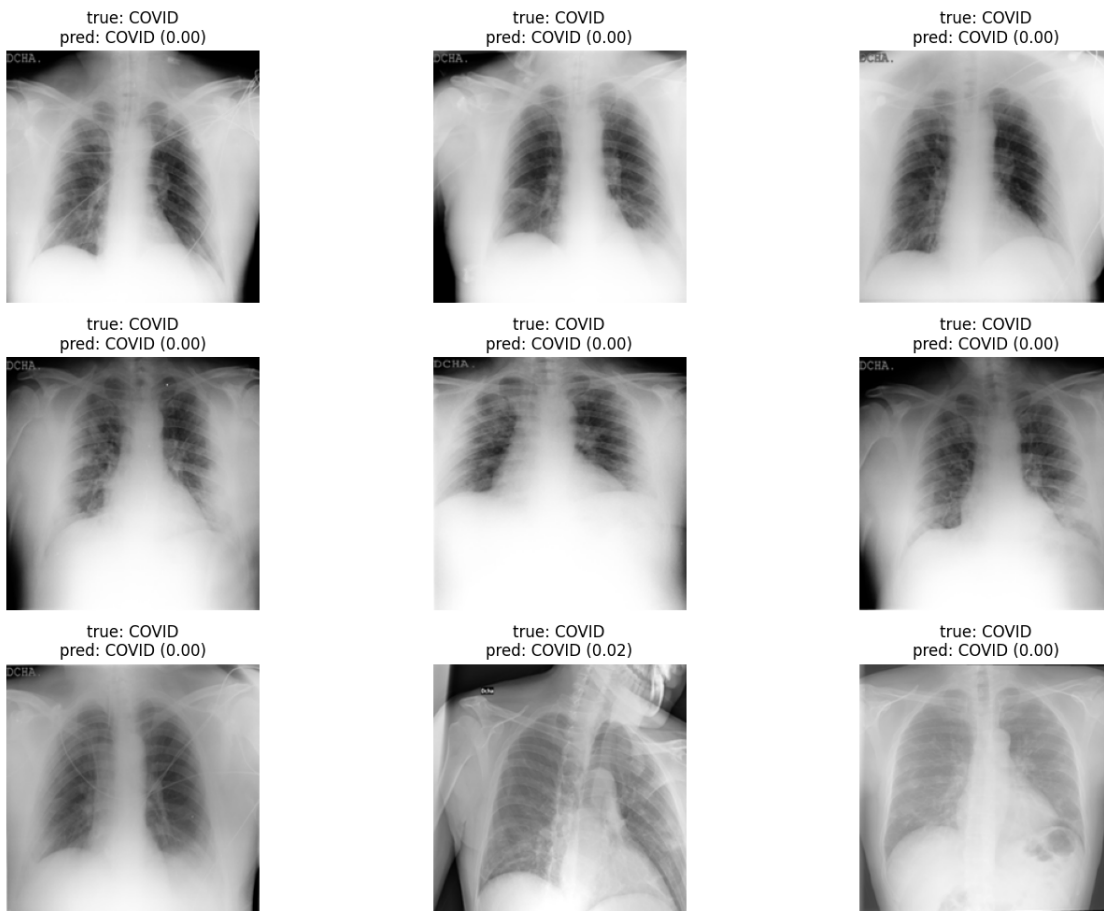
# make predictions
batch_pred_prob = model.predict(images_norm, verbose=0)
batch_pred = (batch_pred_prob > 0.5).astype(int).flatten()

# plot 9 samples
plt.figure(figsize=(15, 10))
for i in range(9):
    plt.subplot(3, 3, i + 1)
```

```
plt.imshow(images_raw[i].astype("uint8"))
true_class = class_names[int(labels_raw[i])]
pred_class = class_names[batch_pred[i]]
prob = batch_pred_prob[i][0]
plt.title(f"true: {true_class}\npred: {pred_class} ({prob:.2f})")
plt.axis("off")

plt.tight_layout()
plt.show()
```

plotting sample predictions with raw images



[ ]: